

# Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

# Importing Datasets

```
In [2]: df=pd.read_csv(r"E:\154\C10_air\csvs_per_year\csvs_per_year\madrid_2004.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	P
0	2004-08-01 01:00:00	NaN	0.66	NaN	NaN	NaN	89.550003	118.900002	NaN	40.020000	39.990000
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.950000
2	2004-08-01 01:00:00	NaN	1.02	NaN	NaN	NaN	93.389999	138.600006	NaN	20.860001	49.480000
3	2004-08-01 01:00:00	NaN	0.53	NaN	NaN	NaN	87.290001	105.000000	NaN	36.730000	31.070000
4	2004-08-01 01:00:00	NaN	0.17	NaN	NaN	NaN	34.910000	35.349998	NaN	86.269997	54.080000
...	...	...	...	...	...	...	...	...	...	...	...
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900000
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	NaN	97.139999	146.899994	2.34	7.740000	37.680000
245493	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.13	102.699997	132.600006	NaN	17.809999	22.840000
245494	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.09	82.599998	102.599998	NaN	NaN	45.630000
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.380000

245496 rows × 17 columns

# Data Cleaning and Data Preprocessing

```
In [3]: df=df.dropna()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      19397 non-null   object 
 1   BEN        19397 non-null   float64
 2   CO         19397 non-null   float64
 3   EBE        19397 non-null   float64
 4   MXY        19397 non-null   float64
 5   NMHC       19397 non-null   float64
 6   NO_2       19397 non-null   float64
 7   NOx        19397 non-null   float64
 8   OXY        19397 non-null   float64
 9   O_3         19397 non-null   float64
 10  PM10       19397 non-null   float64
 11  PM25       19397 non-null   float64
 12  PXY        19397 non-null   float64
 13  SO_2       19397 non-null   float64
 14  TCH         19397 non-null   float64
 15  TOL         19397 non-null   float64
 16  station    19397 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 2.7+ MB
```

```
In [6]: data=df[['CO' , 'station']]  
data
```

Out[6]:

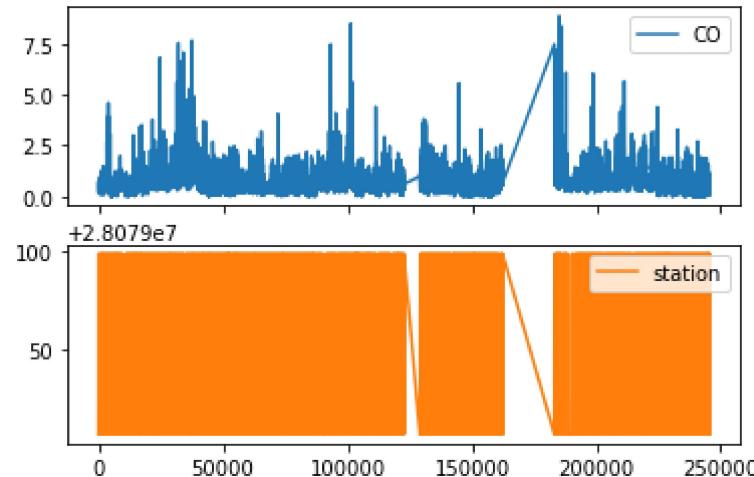
	CO	station
5	0.63	28079006
22	0.36	28079024
26	0.46	28079099
32	0.67	28079006
49	0.30	28079024
...	...	...
245463	0.08	28079024
245467	0.67	28079099
245473	1.12	28079006
245491	0.21	28079024
245495	0.67	28079099

19397 rows × 2 columns

## Line chart

```
In [7]: data.plot.line(subplots=True)
```

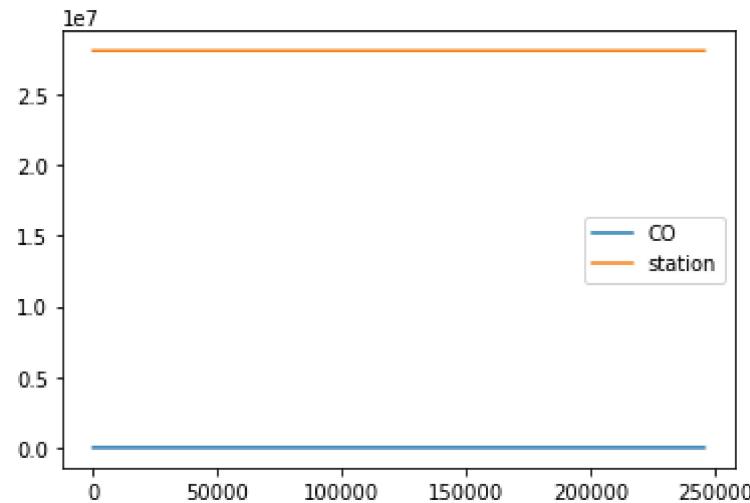
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



## Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

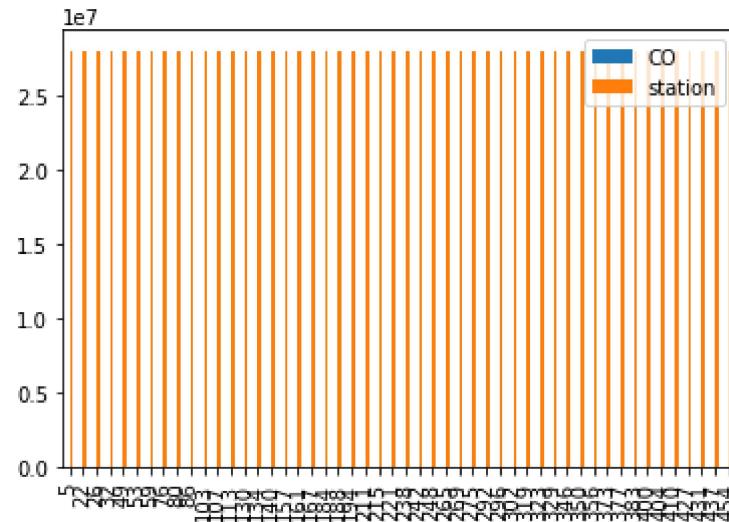


## Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

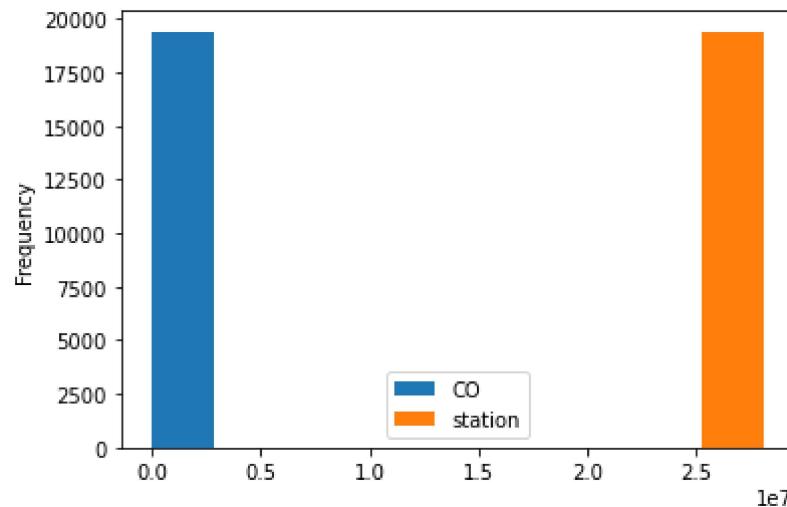
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```

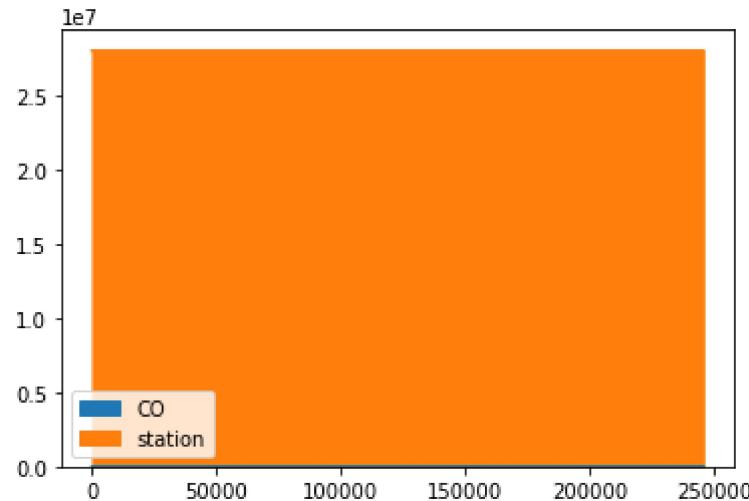
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [12]: data.plot.area()
```

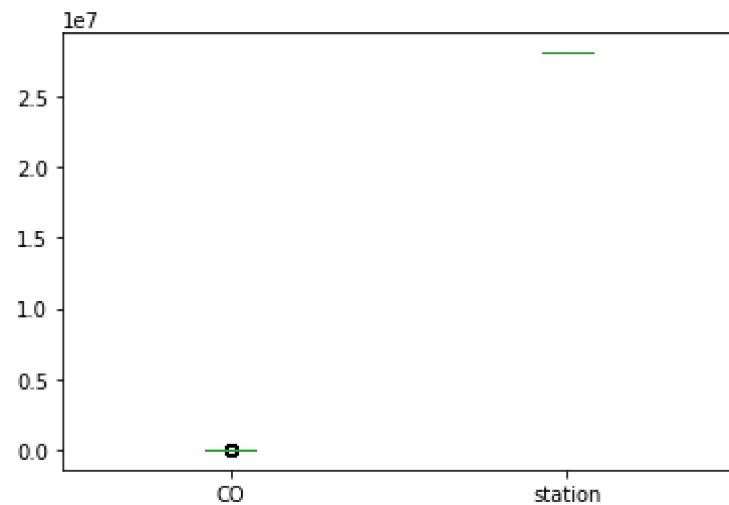
```
Out[12]: <AxesSubplot:>
```



## Box chart

In [13]: `data.plot.box()`

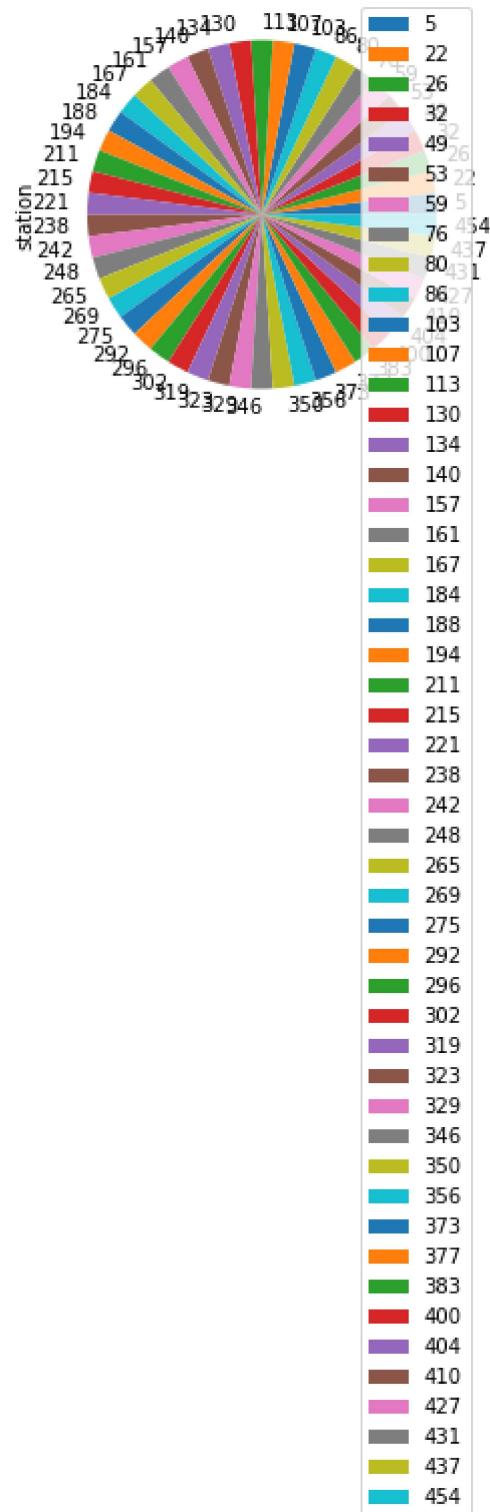
Out[13]: <AxesSubplot:>



## Pie chart

```
In [14]: b.plot.pie(y='station' )
```

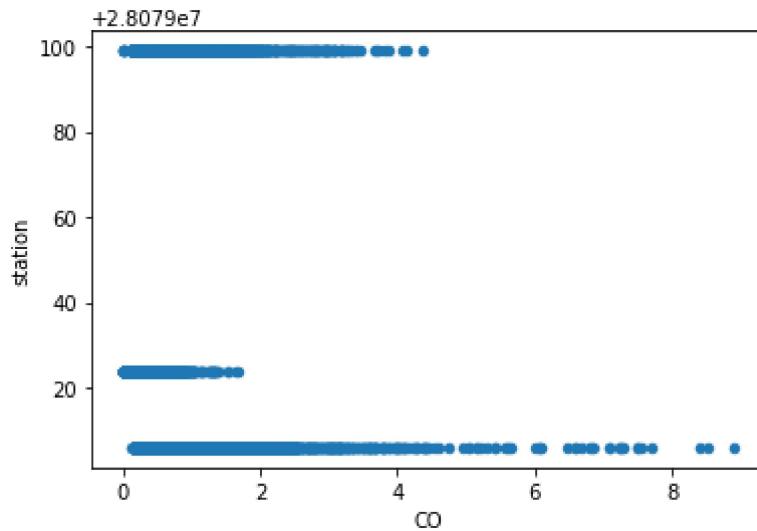
```
Out[14]: <AxesSubplot:ylabel='station'>
```



## Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      19397 non-null   object 
 1   BEN       19397 non-null   float64
 2   CO        19397 non-null   float64
 3   EBE       19397 non-null   float64
 4   MXY       19397 non-null   float64
 5   NMHC      19397 non-null   float64
 6   NO_2      19397 non-null   float64
 7   NOx       19397 non-null   float64
 8   OXY       19397 non-null   float64
 9   O_3        19397 non-null   float64
 10  PM10      19397 non-null   float64
 11  PM25      19397 non-null   float64
 12  PXY       19397 non-null   float64
 13  SO_2      19397 non-null   float64
 14  TCU       19397 non-null   float64
```

```
In [17]: df.describe()
```

Out[17]:

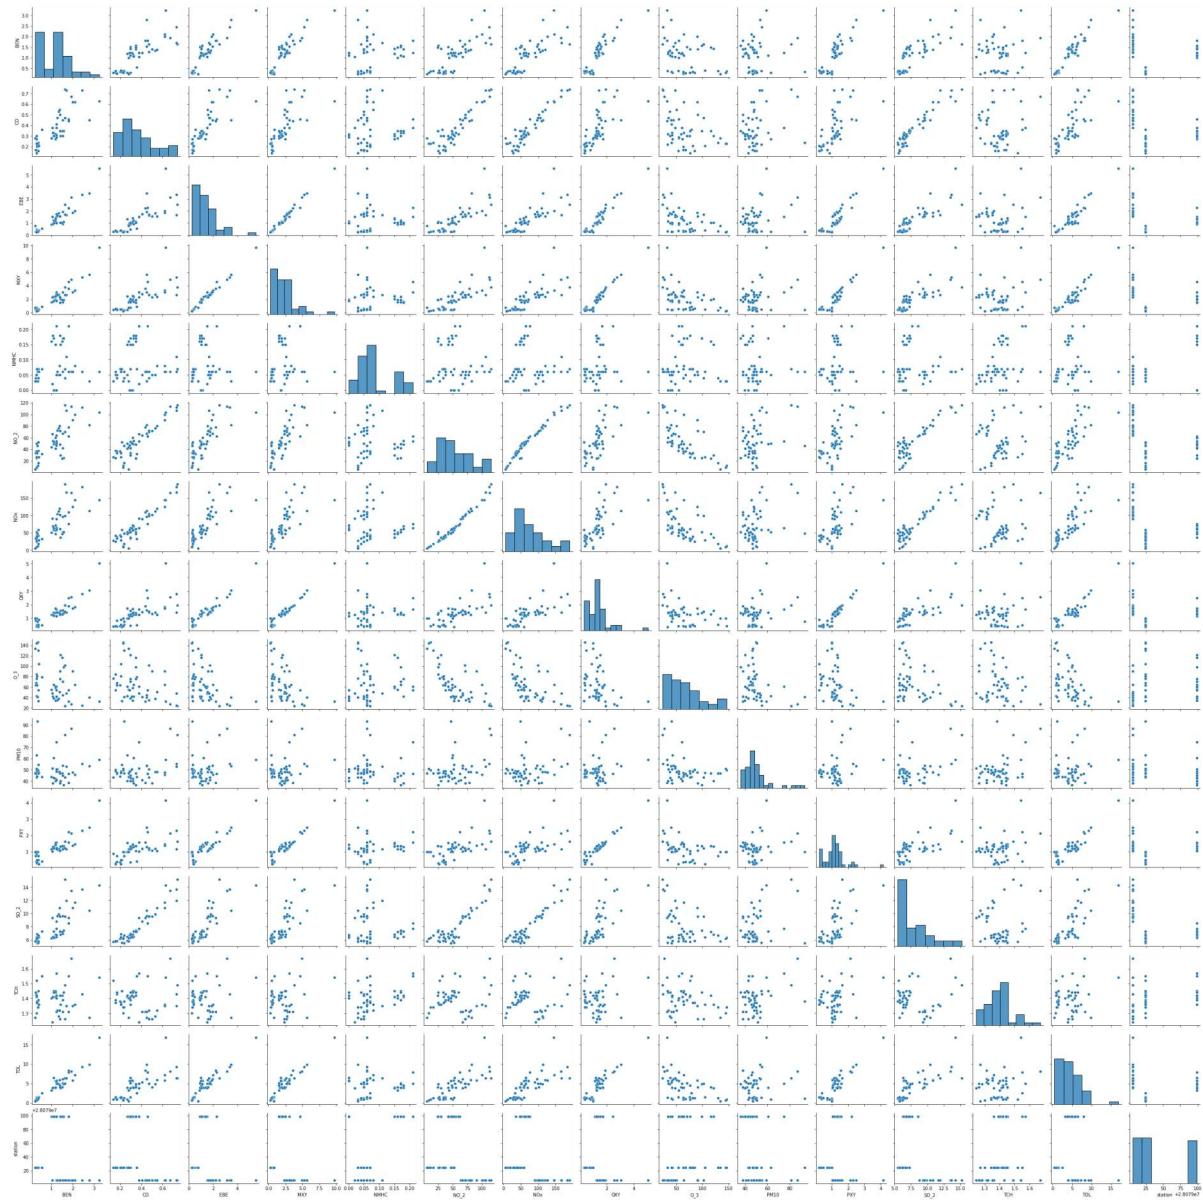
	BEN	CO	EBE	MXY	NMHC	NO_2	
<b>count</b>	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	193
<b>mean</b>	2.250781	0.675347	2.775913	5.424809	0.151024	62.887023	1
<b>std</b>	2.184724	0.591026	2.729622	5.554358	0.158603	37.952255	1
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.090000	
<b>25%</b>	0.870000	0.320000	1.020000	1.780000	0.060000	35.150002	
<b>50%</b>	1.620000	0.520000	1.970000	3.800000	0.110000	58.310001	
<b>75%</b>	2.910000	0.860000	3.580000	7.260000	0.200000	85.730003	1
<b>max</b>	34.180000	8.900000	41.880001	91.599998	4.810000	355.100006	17

```
In [18]: df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

## EDA AND VISUALIZATION

```
In [19]: sns.pairplot(df1[0:50])
```

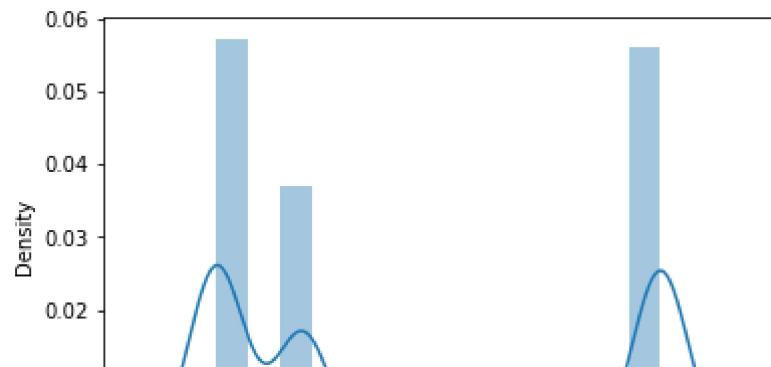
```
Out[19]: <seaborn.axisgrid.PairGrid at 0x239bff23760>
```



In [20]: `sns.distplot(df1['station'])`

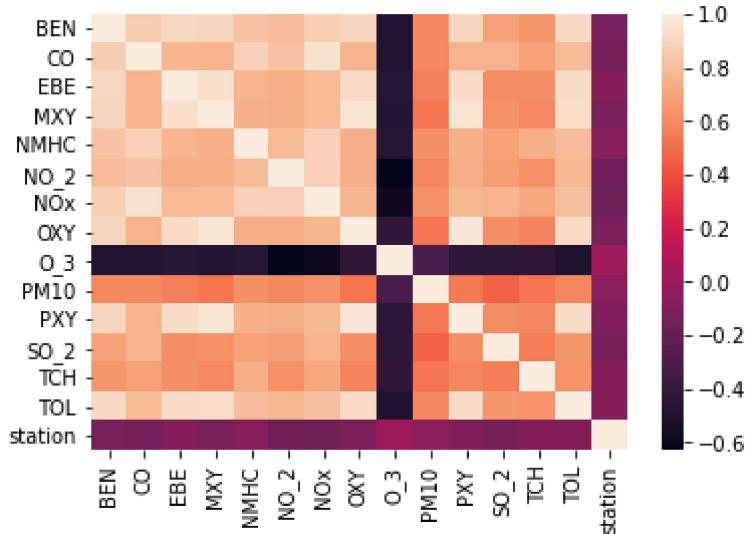
```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <AxesSubplot:>



## TO TRAIN THE MODEL AND MODEL BUILDING

In [22]: `x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['station']`

```
In [23]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: 28079082.290353786
```

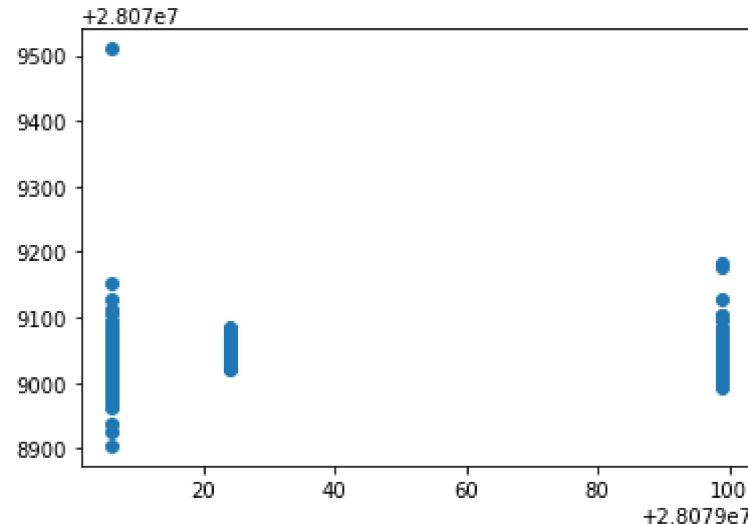
```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[26]:
```

	Co-efficient
BEN	-3.916004
CO	25.450922
EBE	3.588123
MXY	-3.073213
NMHC	113.656409
NO_2	-0.169336
NOx	-0.261784
OXY	-2.255896
O_3	-0.302197
PM10	0.056435
PXY	5.342310
SO_2	-0.288649
TCH	-10.539746
TOL	1.045953

```
In [27]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x239ce9dce20>
```



## ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.08963454278133876
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.11200986318505224
```

## Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.09520468733101284
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.11123283211498791
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

## Accuracy(Lasso)

```
In [35]: la.score(x_train,y_train)
```

```
Out[35]: 0.052392910883565125
```

```
In [36]: la.score(x_test,y_test)
```

```
Out[36]: 0.05312835315738307
```

## ElasticNet

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: en.coef_
```

```
Out[38]: array([-0.          ,  0.41771514,  1.54245331, -1.85100609,  0.          ,
 -0.18707237, -0.07872452, -0.          , -0.22765108,  0.09305377,
 0.29100025, -0.18213256,  0.          ,  1.13534647])
```

```
In [39]: en.intercept_
```

```
Out[39]: 28079068.515236303
```

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

```
Out[41]: 0.0668449030875542
```

## Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

38.42571828936665  
1640.1889361425056  
40.4992461181996

## Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_P10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (19397, 14)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (19397,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: prediction=logr.predict(observation)
```

```
print(prediction)
```

```
[28079006]
```

```
In [52]: logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.7360416559261741
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.9999978255573396
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[9.99997826e-01, 7.75018107e-20, 2.17444266e-06]])
```

## Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
}
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                      param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                  'min_samples_leaf': [5, 10, 15, 20, 25],  
                                  'n_estimators': [10, 20, 30, 40, 50]},  
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.7715253942540938
```

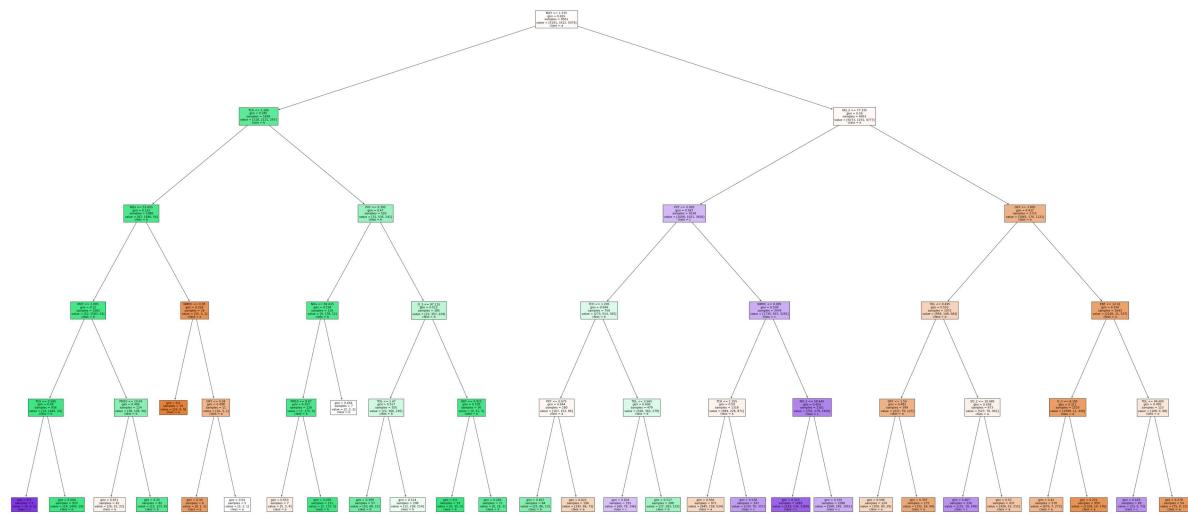
```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b'])
```

```
Out[62]: [Text(2072.5714285714284, 1993.2, 'MXY <= 1.375\ngini = 0.655\nsamples = 8561\nvalue = [5191, 3312, 5074]\nclass = a'),  
Text(956.5714285714284, 1630.8000000000002, 'TCH <= 1.285\ngini = 0.285\nsamples = 1608\nvalue = [118, 2121, 297]\nclass = b'),  
Text(518.1428571428571, 1268.4, 'NOx <= 53.655\ngini = 0.155\nsamples = 1088\nvalue = [87, 1586, 56]\nclass = b'),  
Text(318.85714285714283, 906.0, 'MXY <= 1.005\ngini = 0.12\nsamples = 1062\nvalue = [52, 1583, 54]\nclass = b'),  
Text(159.42857142857142, 543.5999999999999, 'TCH <= 0.585\ngini = 0.05\nsamples = 938\nvalue = [14, 1445, 24]\nclass = b'),  
Text(79.71428571428571, 181.1999999999982, 'gini = 0.0\nsamples = 5\nvalue = [0, 0, 5]\nclass = c'),  
Text(239.1428571428571, 181.1999999999982, 'gini = 0.044\nsamples = 933\nvalue = [14, 1445, 19]\nclass = b'),  
Text(478.2857142857142, 543.5999999999999, 'PM10 <= 10.04\ngini = 0.496\nsamples = 124\nvalue = [38, 138, 30]\nclass = b'),  
Text(398.57142857142856, 181.1999999999982, 'gini = 0.651\nsamples = 42\nvalue = [26, 15, 22]\nclass = a'),  
Text(558.0, 181.1999999999982, 'gini = 0.25\nsamples = 82\nvalue = [12, 123, 8]\nclass = b'),  
Text(717.4285714285713, 906.0, 'NMHC <= 0.05\ngini = 0.226\nsamples = 26\nvalue = [35, 3, 2]\nclass = a'),  
Text(637.7142857142857, 543.5999999999999, 'gini = 0.0\nsamples = 15\nvalue = [25, 0, 0]\nclass = a'),  
Text(797.1428571428571, 543.5999999999999, 'OXY <= 0.58\ngini = 0.498\nsamples = 11\nvalue = [10, 3, 2]\nclass = a'),  
Text(717.4285714285713, 181.1999999999982, 'gini = 0.34\nsamples = 6\nvalue = [8, 1, 1]\nclass = a'),  
Text(876.8571428571428, 181.1999999999982, 'gini = 0.64\nsamples = 5\nvalue = [2, 2, 1]\nclass = a'),  
Text(1395.0, 1268.4, 'PXY <= 0.395\ngini = 0.47\nsamples = 520\nvalue = [31, 535, 241]\nclass = b'),  
Text(1195.7142857142856, 906.0, 'NOx <= 82.615\ngini = 0.194\nsamples = 135\nvalue = [9, 178, 12]\nclass = b'),  
Text(1116.0, 543.5999999999999, 'PM10 <= 8.87\ngini = 0.157\nsamples = 128\nvalue = [7, 175, 9]\nclass = b'),  
Text(1036.2857142857142, 181.1999999999982, 'gini = 0.653\nsamples = 7\nvalue = [5, 3, 4]\nclass = a'),  
Text(1195.7142857142856, 181.1999999999982, 'gini = 0.076\nsamples = 121\nvalue = [2, 172, 5]\nclass = b'),  
Text(1275.4285714285713, 543.5999999999999, 'gini = 0.656\nsamples = 7\nvalue = [2, 3, 3]\nclass = b'),  
Text(1594.2857142857142, 906.0, 'O_3 <= 97.115\ngini = 0.512\nsamples = 385\nvalue = [22, 357, 229]\nclass = b'),  
Text(1434.8571428571427, 543.5999999999999, 'TOL <= 1.47\ngini = 0.527\nsamples = 355\nvalue = [22, 306, 226]\nclass = b'),  
Text(1355.142857142857, 181.1999999999982, 'gini = 0.399\nsamples = 57\nvalue = [10, 68, 12]\nclass = b'),  
Text(1514.5714285714284, 181.1999999999982, 'gini = 0.524\nsamples = 298\nvalue = [12, 238, 214]\nclass = b'),  
Text(1753.7142857142856, 543.5999999999999, 'OXY <= 0.915\ngini = 0.105\nsamples = 30\nvalue = [0, 51, 3]\nclass = b'),  
Text(1673.9999999999998, 181.1999999999982, 'gini = 0.0\nsamples = 19\nvalue = [0, 35, 0]\nclass = b'),  
Text(1833.4285714285713, 181.1999999999982, 'gini = 0.266\nsamples = 11\nvalue = [0, 16, 3]\nclass = b'),  
Text(3188.5714285714284, 1630.8000000000002, 'NO_2 <= 77.335\ngini = 0.59\nsamples = 1000\nvalue = [1000, 0, 0]\nclass = b')]
```

```
amples = 6953\nvalue = [5073, 1191, 4777]\nclass = a'),  
Text(2550.8571428571427, 1268.4, 'PXY <= 0.905\ngini = 0.587\nsamples = 4238  
\nvalue = [2008, 1021, 3656]\nclass = c'),  
Text(2232.0, 906.0, 'TCH <= 1.295\ngini = 0.644\nsamples = 759\nvalue = [27  
3, 514, 365]\nclass = b'),  
Text(2072.5714285714284, 543.5999999999999, 'PXY <= 0.675\ngini = 0.644\nsam  
ples = 280\nvalue = [167, 152, 86]\nclass = a'),  
Text(1992.8571428571427, 181.1999999999982, 'gini = 0.467\nsamples = 84\nva  
lue = [25, 86, 13]\nclass = b'),  
Text(2152.285714285714, 181.1999999999982, 'gini = 0.622\nsamples = 196\nva  
lue = [142, 66, 73]\nclass = a'),  
Text(2391.428571428571, 543.5999999999999, 'TOL <= 3.565\ngini = 0.606\nsam  
ples = 479\nvalue = [106, 362, 279]\nclass = b'),  
Text(2311.7142857142853, 181.1999999999982, 'gini = 0.626\nsamples = 191\nva  
lue = [69, 79, 146]\nclass = c'),  
Text(2471.142857142857, 181.1999999999982, 'gini = 0.517\nsamples = 288\nva  
lue = [37, 283, 133]\nclass = b'),  
Text(2869.7142857142853, 906.0, 'NMHC <= 0.095\ngini = 0.539\nsamples = 3479  
\nvalue = [1735, 507, 3291]\nclass = c'),  
Text(2710.285714285714, 543.5999999999999, 'TCH <= 1.355\ngini = 0.59\nsampl  
es = 1318\nvalue = [984, 228, 871]\nclass = a'),  
Text(2630.5714285714284, 181.1999999999982, 'gini = 0.566\nsamples = 971\nva  
lue = [845, 158, 534]\nclass = a'),  
Text(2790.0, 181.1999999999982, 'gini = 0.538\nsamples = 347\nvalue = [139,  
70, 337]\nclass = c'),  
Text(3029.142857142857, 543.5999999999999, 'NO_2 <= 58.645\ngini = 0.454\nsa  
mples = 2161\nvalue = [751, 279, 2420]\nclass = c'),  
Text(2949.428571428571, 181.1999999999982, 'gini = 0.323\nsamples = 1065\nva  
lue = [183, 134, 1369]\nclass = c'),  
Text(3108.8571428571427, 181.1999999999982, 'gini = 0.535\nsamples = 1096\nva  
lue = [568, 145, 1051]\nclass = c'),  
Text(3826.2857142857138, 1268.4, 'OXY <= 3.805\ngini = 0.437\nsamples = 2715  
\nvalue = [3065, 170, 1121]\nclass = a'),  
Text(3507.428571428571, 906.0, 'TOL <= 9.495\ngini = 0.552\nsamples = 1072\nva  
lue = [959, 149, 584]\nclass = a'),  
Text(3347.999999999995, 543.5999999999999, 'OXY <= 1.54\ngini = 0.483\nsam  
ples = 399\nvalue = [432, 79, 123]\nclass = a'),  
Text(3268.285714285714, 181.1999999999982, 'gini = 0.596\nsamples = 124\nva  
lue = [100, 60, 29]\nclass = a'),  
Text(3427.7142857142853, 181.1999999999982, 'gini = 0.397\nsamples = 275\nva  
lue = [332, 19, 94]\nclass = a'),  
Text(3666.8571428571427, 543.5999999999999, 'SO_2 <= 10.685\ngini = 0.558\nsa  
mples = 673\nvalue = [527, 70, 461]\nclass = a'),  
Text(3587.142857142857, 181.1999999999982, 'gini = 0.467\nsamples = 236\nva  
lue = [101, 18, 246]\nclass = c'),  
Text(3746.5714285714284, 181.1999999999982, 'gini = 0.52\nsamples = 437\nva  
lue = [426, 52, 215]\nclass = a'),  
Text(4145.142857142857, 906.0, 'EBE <= 12.01\ngini = 0.334\nsamples = 1643\nva  
lue = [2106, 21, 537]\nclass = a'),  
Text(3985.7142857142853, 543.5999999999999, 'O_3 <= 8.185\ngini = 0.311\nsam  
ples = 1520\nvalue = [1998, 21, 449]\nclass = a'),  
Text(3905.999999999995, 181.1999999999982, 'gini = 0.42\nsamples = 570\nva  
lue = [670, 7, 273]\nclass = a'),  
Text(4065.428571428571, 181.1999999999982, 'gini = 0.221\nsamples = 950\nva  
lue = [1328, 14, 176]\nclass = a'),  
Text(4304.571428571428, 543.5999999999999, 'TOL <= 49.605\ngini = 0.495\nsa  
mples = 123\nvalue = [108, 0, 88]\nclass = a'),
```

```
Text(4224.857142857142, 181.19999999999982, 'gini = 0.429\nsamples = 69\nvalue = [33, 0, 73]\nclass = c'),  
Text(4384.285714285714, 181.19999999999982, 'gini = 0.278\nsamples = 54\nvalue = [75, 0, 15]\nclass = a')]
```



## Accuracy

### Linear Regression

```
In [63]: lr.score(x_train,y_train)
```

```
Out[63]: 0.11200986318505224
```

### Ridge Regression

```
In [64]: rr.score(x_train,y_train)
```

```
Out[64]: 0.11123283211498791
```

### Lasso Regression

```
In [65]: la.score(x_test,y_test)
```

```
Out[65]: 0.05312835315738307
```

## ElasticNet Regression

```
In [66]: en.score(x_test,y_test)
```

```
Out[66]: 0.0668449030875542
```

## Logistic Regression

```
In [67]: logr.score(fs,target_vector)
```

```
Out[67]: 0.7360416559261741
```

## Random Forest

```
In [68]: grid_search.best_score_
```

```
Out[68]: 0.7715253942540938
```

## Conclusion

**Random Forest is suitable for this dataset**