

# Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

# Importing Datasets

```
In [2]: df=pd.read_csv(r"E:\154\C10_air\csvs_per_year\csvs_per_year\madrid_2005.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10
0	2005-11-01 01:00:00	NaN	0.77	NaN	NaN	NaN	57.130001	128.699997	NaN	14.720000	14.91
1	2005-11-01 01:00:00	1.52	0.65	1.49	4.57	0.25	86.559998	181.699997	1.27	11.680000	30.93
2	2005-11-01 01:00:00	NaN	0.40	NaN	NaN	NaN	46.119999	53.000000	NaN	30.469999	14.60
3	2005-11-01 01:00:00	NaN	0.42	NaN	NaN	NaN	37.220001	52.009998	NaN	21.379999	15.16
4	2005-11-01 01:00:00	NaN	0.57	NaN	NaN	NaN	32.160000	36.680000	NaN	33.410000	5.00
...	...	...	...	...	...	...	...	...	...	...	...
236995	2006-01-01 00:00:00	1.08	0.36	1.01	NaN	0.11	21.990000	23.610001	NaN	43.349998	5.00
236996	2006-01-01 00:00:00	0.39	0.54	1.00	1.00	0.11	2.200000	4.220000	1.00	69.639999	4.95
236997	2006-01-01 00:00:00	0.19	NaN	0.26	NaN	0.08	26.730000	30.809999	NaN	43.840000	4.31
236998	2006-01-01 00:00:00	0.14	NaN	1.00	NaN	0.06	13.770000	17.770000	NaN	NaN	5.00
236999	2006-01-01 00:00:00	0.50	0.40	0.73	1.84	0.13	20.940001	26.950001	1.49	48.259998	5.67

237000 rows × 17 columns

# Data Cleaning and Data Preprocessing

```
In [3]: df=df.dropna()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20070 entries, 5 to 236999
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      20070 non-null   object 
 1   BEN        20070 non-null   float64
 2   CO         20070 non-null   float64
 3   EBE        20070 non-null   float64
 4   MXY        20070 non-null   float64
 5   NMHC       20070 non-null   float64
 6   NO_2       20070 non-null   float64
 7   NOx        20070 non-null   float64
 8   OXY        20070 non-null   float64
 9   O_3         20070 non-null   float64
 10  PM10       20070 non-null   float64
 11  PM25       20070 non-null   float64
 12  PXY        20070 non-null   float64
 13  SO_2       20070 non-null   float64
 14  TCH         20070 non-null   float64
 15  TOL         20070 non-null   float64
 16  station    20070 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 2.8+ MB
```

```
In [6]: data=df[['CO' , 'station']]  
data
```

Out[6]:

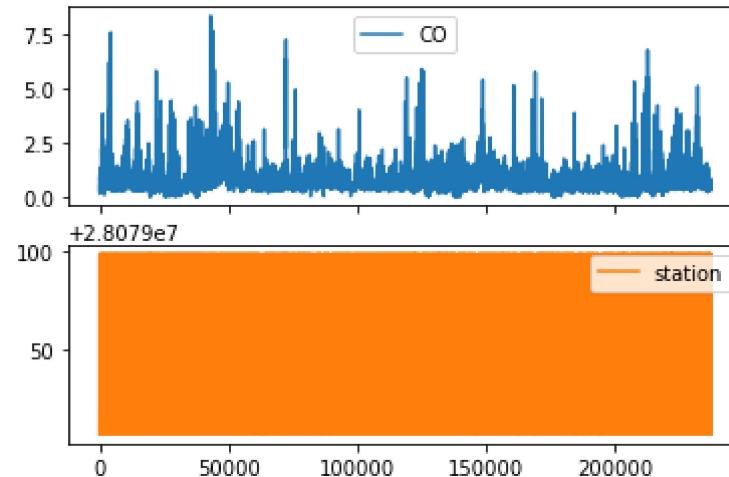
	CO	station
5	0.88	28079006
22	0.22	28079024
25	0.49	28079099
31	0.84	28079006
48	0.20	28079024
...	...	...
236970	0.39	28079024
236973	0.45	28079099
236979	0.38	28079006
236996	0.54	28079024
236999	0.40	28079099

20070 rows × 2 columns

## Line chart

```
In [7]: data.plot.line(subplots=True)
```

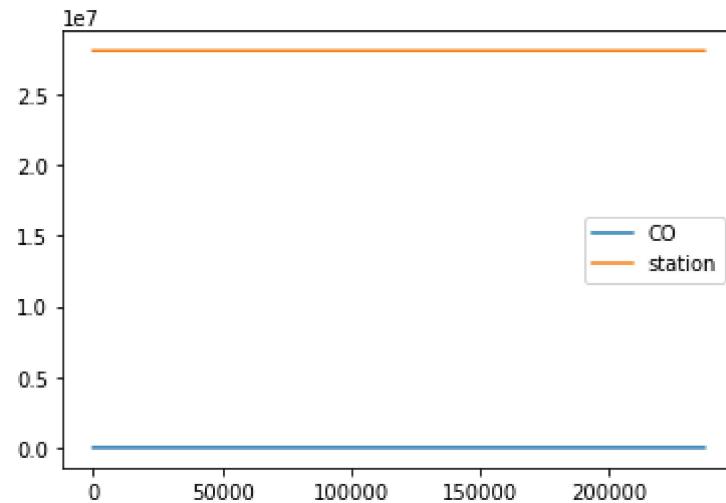
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



## Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

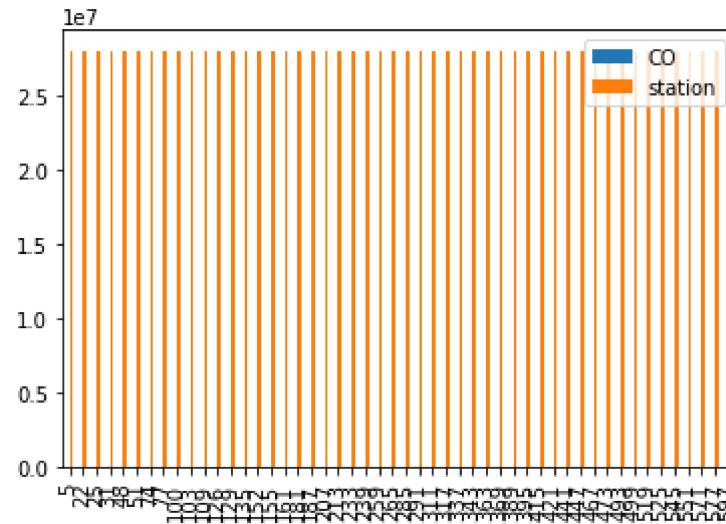


## Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

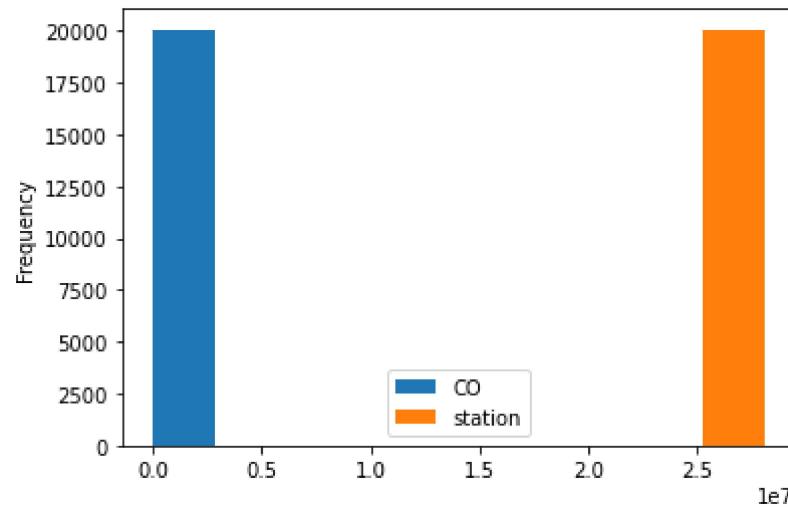
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```

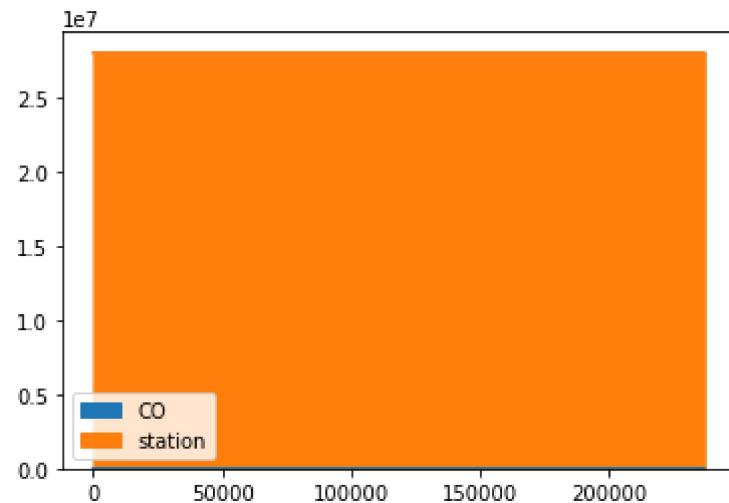
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [12]: data.plot.area()
```

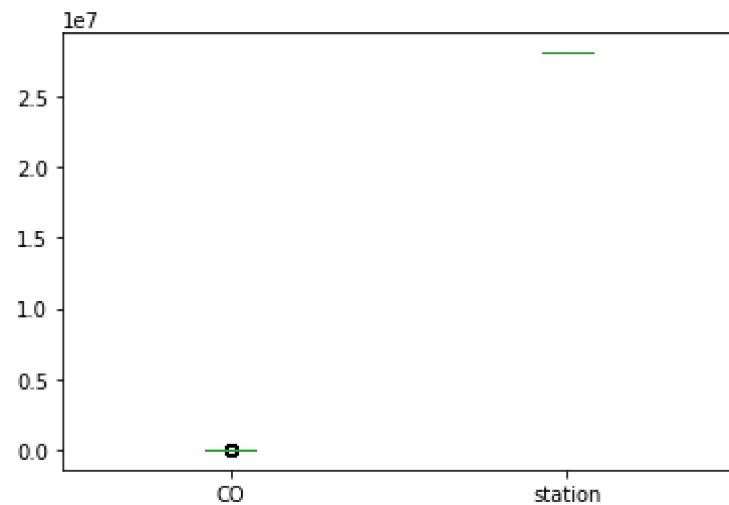
```
Out[12]: <AxesSubplot:>
```



## Box chart

In [13]: `data.plot.box()`

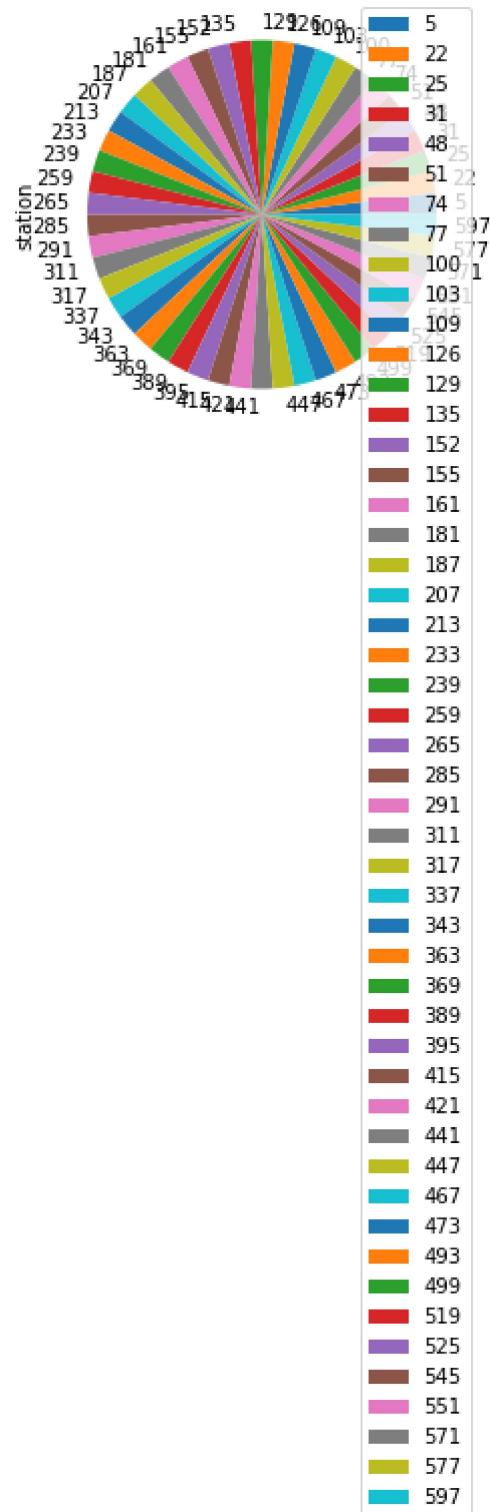
Out[13]: <AxesSubplot:>



## Pie chart

```
In [14]: b.plot.pie(y='station' )
```

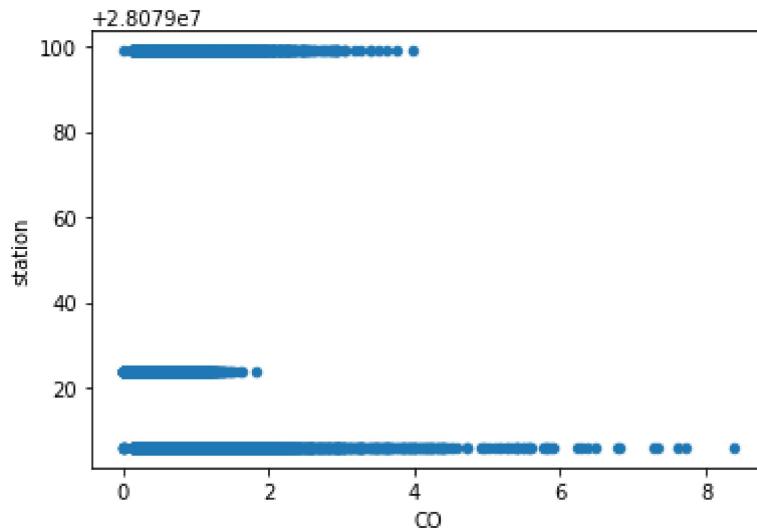
```
Out[14]: <AxesSubplot:ylabel='station'>
```



## Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20070 entries, 5 to 236999
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   date      20070 non-null   object 
 1   BEN       20070 non-null   float64
 2   CO        20070 non-null   float64
 3   EBE       20070 non-null   float64
 4   MXY       20070 non-null   float64
 5   NMHC      20070 non-null   float64
 6   NO_2      20070 non-null   float64
 7   NOx       20070 non-null   float64
 8   OXY       20070 non-null   float64
 9   O_3        20070 non-null   float64
 10  PM10      20070 non-null   float64
 11  PM25      20070 non-null   float64
 12  PXY       20070 non-null   float64
 13  SO_2      20070 non-null   float64
 14  TCU       20070 non-null   float64
```

In [17]: `df.describe()`

Out[17]:

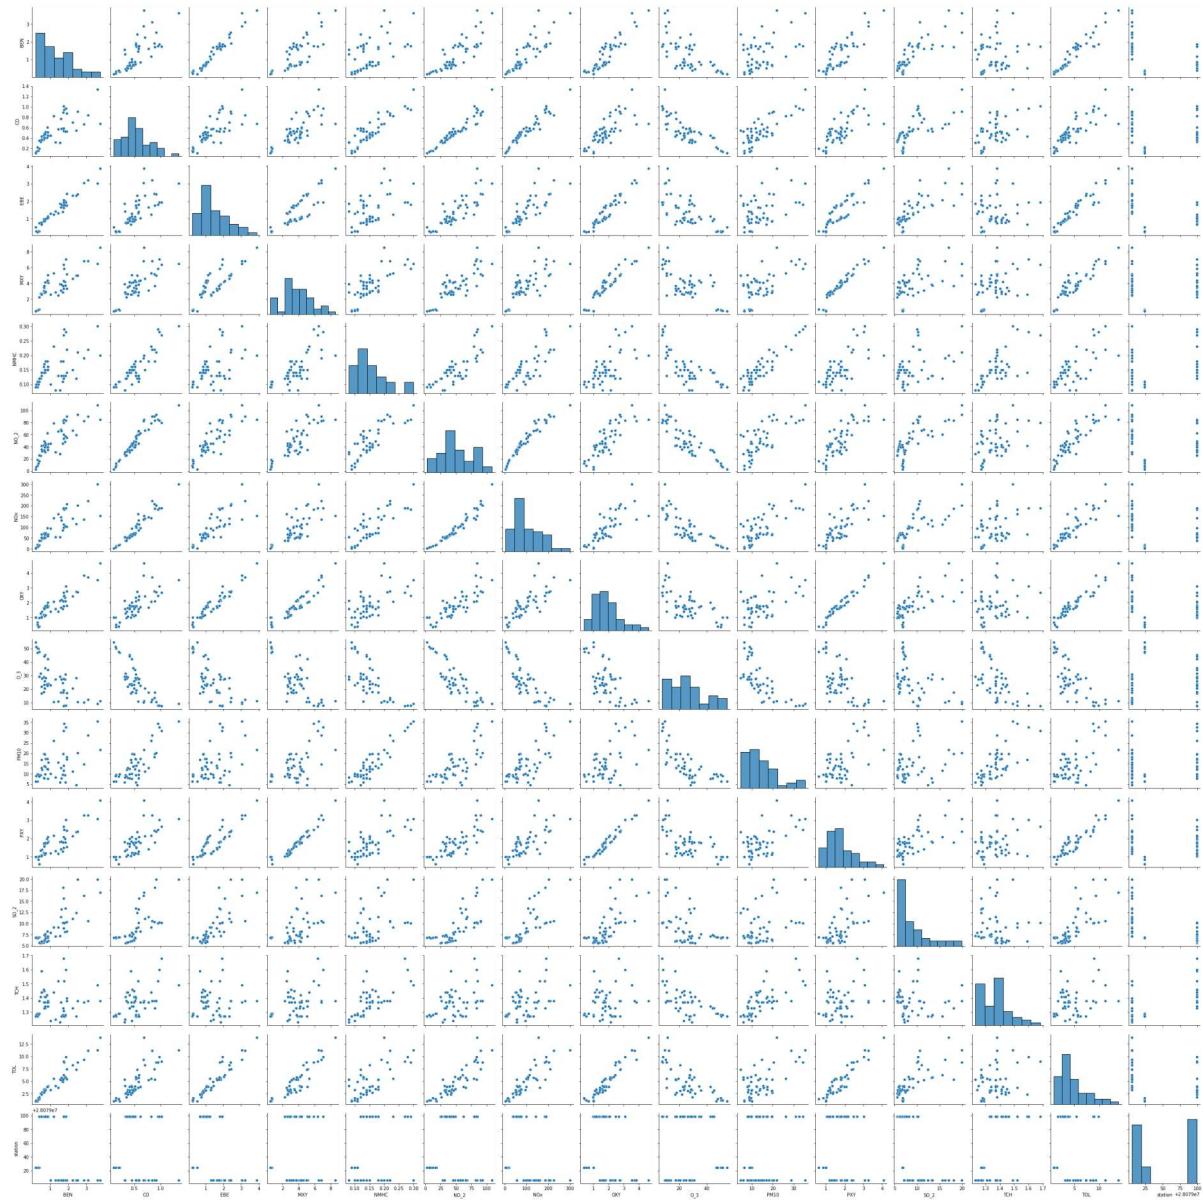
	BEN	CO	EBE	MXY	NMHC	NO_2	
<b>count</b>	20070.000000	20070.000000	20070.000000	20070.000000	20070.000000	20070.000000	200
<b>mean</b>	1.923656	0.720657	2.345423	5.457855	0.179282	66.226924	1
<b>std</b>	2.019061	0.549723	2.379219	5.495147	0.152783	40.568197	1
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	0.690000	0.400000	0.950000	1.930000	0.090000	36.602499	
<b>50%</b>	1.260000	0.580000	1.480000	3.800000	0.150000	60.525000	1
<b>75%</b>	2.510000	0.880000	2.950000	7.210000	0.220000	89.317499	1
<b>max</b>	26.570000	8.380000	29.870001	71.050003	1.880000	419.500000	17

In [18]: `df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`

## EDA AND VISUALIZATION

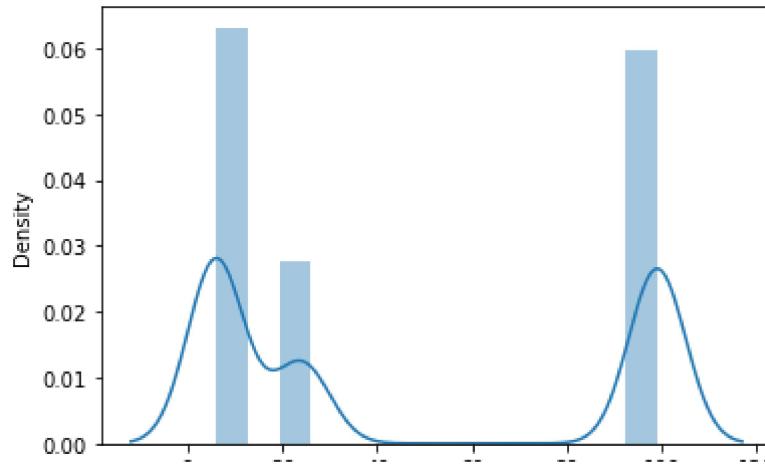
```
In [19]: sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x2cdeb13820>
```



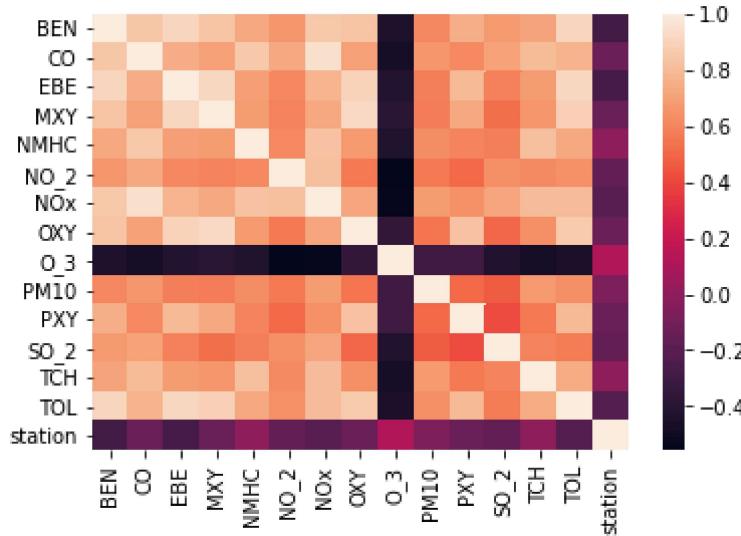
```
In [20]: sns.distplot(df1['station'])
FutureWarning: Please adapt your code to use either distplot (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [23]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: 28078951.808543332
```

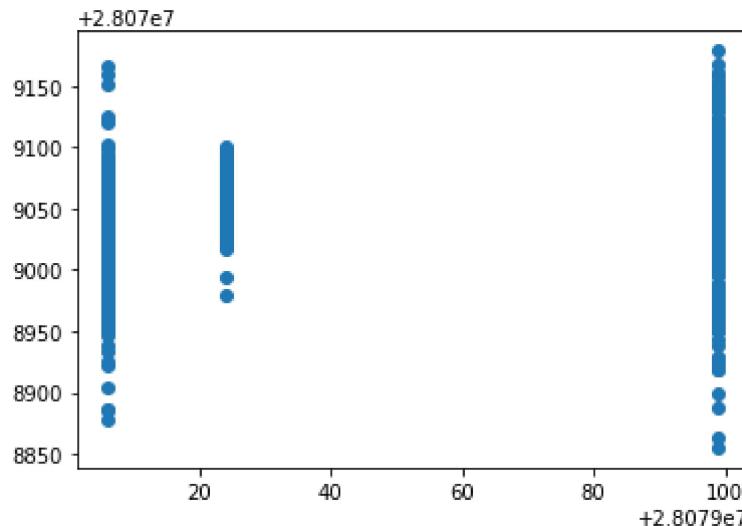
```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[26]:
```

	Co-efficient
BEN	-9.465611
CO	38.269204
EBE	-13.753835
MXY	3.609269
NMHC	78.684788
NO_2	0.121025
NOx	-0.264759
OXY	3.559233
O_3	0.023548
PM10	0.023321
PXY	2.566572
SO_2	0.185111
TCH	68.698585
TOL	-0.566203

```
In [27]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x2cdafa882250>
```



## ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.28233362743052093
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.31282390667177273
```

## Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.28275844440462194
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.3125741505664318
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

## Accuracy(Lasso)

```
In [35]: la.score(x_train,y_train)
```

```
Out[35]: 0.06293246283731035
```

```
In [36]: la.score(x_test,y_test)
```

```
Out[36]: 0.06685746388029112
```

## ElasticNet

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: en.coef_
```

```
Out[38]: array([-5.73672840e+00,  1.52481358e+00, -7.59109090e+00,  2.67067424e+00,
   8.76231211e-01, -6.12971883e-02, -2.17032038e-03,  2.00575865e+00,
  -1.84742649e-02,  2.14922961e-01,  1.42783112e+00,  1.27260444e-01,
  1.60778992e+00, -7.90151637e-01])
```

```
In [39]: en.intercept_
```

```
Out[39]: 28079050.583618924
```

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

```
Out[41]: 0.17339846972558737
```

## Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

36.81280710465069  
1534.686303051333  
39.175072470275445

## Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_P10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (20070, 14)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (20070,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: prediction=logr.predict(observation)
```

```
print(prediction)
```

```
[28079006]
```

```
In [52]: logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.879023418036871
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.9998967601812779
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[9.99896760e-01, 3.21124597e-30, 1.03239819e-04]])
```

## Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
}
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                      param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                  'min_samples_leaf': [5, 10, 15, 20, 25],  
                                  'n_estimators': [10, 20, 30, 40, 50]},  
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

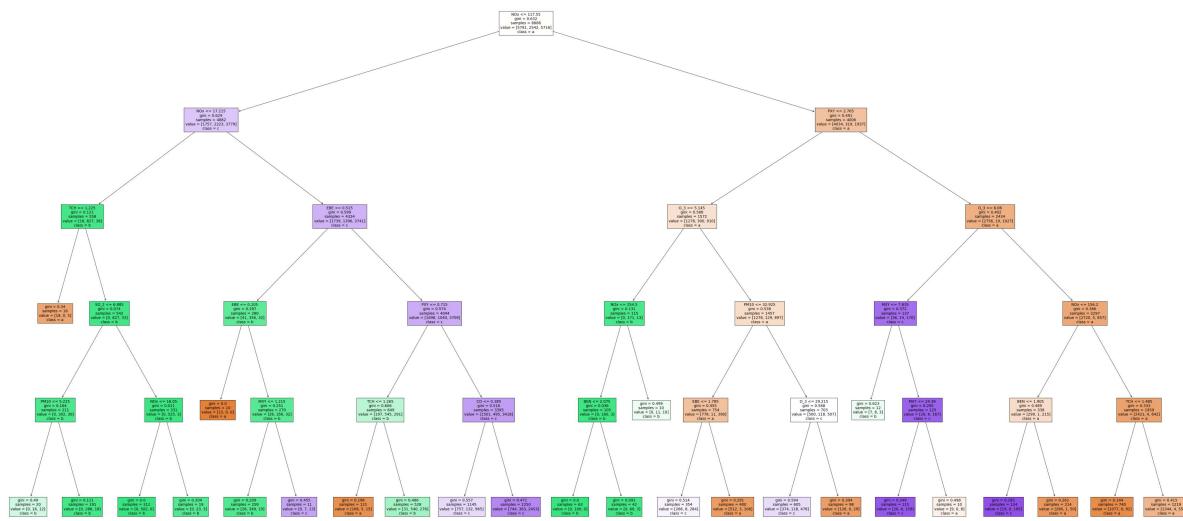
```
Out[60]: 0.8633356706847494
```

```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b'])
```

```
Out[62]: [Text(1965.6818181818182, 1993.2, 'NOx <= 117.55\ngini = 0.632\nsamples = 888\nvalue = [5791, 2542, 5716]\nnclass = a'),  
 Text(786.2727272727273, 1630.8000000000002, 'NOx <= 17.215\ngini = 0.629\nsamples = 4882\nvalue = [1757, 2223, 3779]\nnclass = c'),  
 Text(304.3636363636364, 1268.4, 'TCH <= 1.225\ngini = 0.121\nsamples = 558\nvalue = [18, 827, 38]\nnclass = b'),  
 Text(202.9090909090909, 906.0, 'gini = 0.34\nsamples = 16\nvalue = [18, 0, 5]\nnclass = a'),  
 Text(405.8181818181818, 906.0, 'SO_2 <= 6.885\ngini = 0.074\nsamples = 542\nvalue = [0, 827, 33]\nnclass = b'),  
 Text(202.9090909090909, 543.5999999999999, 'PM10 <= 5.225\ngini = 0.164\nsamples = 211\nvalue = [0, 302, 30]\nnclass = b'),  
 Text(101.45454545454545, 181.1999999999982, 'gini = 0.49\nsamples = 20\nvalue = [0, 16, 12]\nnclass = b'),  
 Text(304.3636363636364, 181.1999999999982, 'gini = 0.111\nsamples = 191\nvalue = [0, 286, 18]\nnclass = b'),  
 Text(608.7272727272727, 543.5999999999999, 'NOx <= 16.05\ngini = 0.011\nsamples = 331\nvalue = [0, 525, 3]\nnclass = b'),  
 Text(507.27272727272725, 181.1999999999982, 'gini = 0.0\nsamples = 312\nvalue = [0, 502, 0]\nnclass = b'),  
 Text(710.18181818181, 181.1999999999982, 'gini = 0.204\nsamples = 19\nvalue = [0, 23, 3]\nnclass = b'),  
 Text(1268.18181818182, 1268.4, 'EBE <= 0.515\ngini = 0.599\nsamples = 4324\nvalue = [1739, 1396, 3741]\nnclass = c'),  
 Text(913.0909090909091, 906.0, 'EBE <= 0.105\ngini = 0.297\nsamples = 280\nvalue = [41, 356, 32]\nnclass = b'),  
 Text(811.6363636363636, 543.5999999999999, 'gini = 0.0\nsamples = 10\nvalue = [15, 0, 0]\nnclass = a'),  
 Text(1014.5454545454545, 543.5999999999999, 'MXY <= 1.215\ngini = 0.251\nsamples = 270\nvalue = [26, 356, 32]\nnclass = b'),  
 Text(913.0909090909091, 181.1999999999982, 'gini = 0.209\nsamples = 259\nvalue = [26, 349, 19]\nnclass = b'),  
 Text(1116.0, 181.1999999999982, 'gini = 0.455\nsamples = 11\nvalue = [0, 7, 13]\nnclass = c'),  
 Text(1623.2727272727273, 906.0, 'PXY <= 0.715\ngini = 0.574\nsamples = 4044\nvalue = [1698, 1040, 3709]\nnclass = c'),  
 Text(1420.3636363636363, 543.5999999999999, 'TCH <= 1.265\ngini = 0.606\nsamples = 649\nvalue = [197, 545, 291]\nnclass = b'),  
 Text(1318.909090909091, 181.1999999999982, 'gini = 0.196\nsamples = 121\nvalue = [166, 5, 15]\nnclass = a'),  
 Text(1521.81818181818, 181.1999999999982, 'gini = 0.486\nsamples = 528\nvalue = [31, 540, 276]\nnclass = b'),  
 Text(1826.18181818182, 543.5999999999999, 'CO <= 0.385\ngini = 0.516\nsamples = 3395\nvalue = [1501, 495, 3418]\nnclass = c'),  
 Text(1724.7272727272727, 181.1999999999982, 'gini = 0.557\nsamples = 1145\nvalue = [757, 132, 965]\nnclass = c'),  
 Text(1927.6363636363635, 181.1999999999982, 'gini = 0.471\nsamples = 2250\nvalue = [744, 363, 2453]\nnclass = c'),  
 Text(3145.090909090909, 1630.8000000000002, 'PXY <= 2.765\ngini = 0.491\nsamples = 4006\nvalue = [4034, 319, 1937]\nnclass = a'),  
 Text(2587.090909090909, 1268.4, 'O_3 <= 5.145\ngini = 0.588\nsamples = 1572\nvalue = [1278, 300, 910]\nnclass = a'),  
 Text(2333.4545454545455, 906.0, 'NOx <= 254.5\ngini = 0.131\nsamples = 115\nvalue = [0, 171, 13]\nnclass = b'),  
 Text(2232.0, 543.5999999999999, 'BEN <= 2.075\ngini = 0.036\nsamples = 105\nvalue = [0, 160, 3]\nnclass = b'),  
 Text(2130.5454545454545, 181.1999999999982, 'gini = 0.0\nsamples = 63\nvalue = [0, 159, 1]\nnclass = a')]
```

```
e = [0, 100, 0]\nclass = b'),  
    Text(2333.4545454545455, 181.19999999999982, 'gini = 0.091\nclass = 42\nvalue = [0, 60, 3]\nclass = b'),  
    Text(2434.909090909091, 543.5999999999999, 'gini = 0.499\nclass = 10\nvalue = [0, 11, 10]\nclass = b'),  
    Text(2840.72727272725, 906.0, 'PM10 <= 32.925\ngini = 0.538\nclass = 145\nvalue = [1278, 129, 897]\nclass = a'),  
    Text(2637.8181818182, 543.5999999999999, 'EBE <= 1.785\ngini = 0.455\nclasses = 754\nvalue = [778, 11, 390]\nclass = a'),  
    Text(2536.3636363636365, 181.19999999999982, 'gini = 0.514\nclass = 354\nvalue = [266, 8, 284]\nclass = c'),  
    Text(2739.272727272727, 181.19999999999982, 'gini = 0.291\nclass = 400\nvalue = [512, 3, 106]\nclass = a'),  
    Text(3043.6363636363635, 543.5999999999999, 'O_3 <= 29.215\ngini = 0.588\nclasses = 703\nvalue = [500, 118, 507]\nclass = c'),  
    Text(2942.181818181818, 181.19999999999982, 'gini = 0.594\nclass = 605\nvalue = [374, 118, 478]\nclass = c'),  
    Text(3145.090909090909, 181.19999999999982, 'gini = 0.304\nclass = 98\nvalue = [126, 0, 29]\nclass = a'),  
    Text(3703.090909090909, 1268.4, 'O_3 <= 6.06\ngini = 0.402\nclass = 2434\nvalue = [2756, 19, 1027]\nclass = a'),  
    Text(3348.0, 906.0, 'MXY <= 7.835\ngini = 0.372\nclass = 137\nvalue = [36, 14, 170]\nclass = c'),  
    Text(3246.5454545454545, 543.5999999999999, 'gini = 0.623\nclass = 12\nvalue = [7, 8, 3]\nclass = b'),  
    Text(3449.4545454545455, 543.5999999999999, 'MXY <= 24.39\ngini = 0.295\nclasses = 125\nvalue = [29, 6, 167]\nclass = c'),  
    Text(3348.0, 181.19999999999982, 'gini = 0.249\nclass = 115\nvalue = [20, 6, 159]\nclass = c'),  
    Text(3550.909090909091, 181.19999999999982, 'gini = 0.498\nclass = 10\nvalue = [9, 0, 8]\nclass = a'),  
    Text(4058.1818181818, 906.0, 'NOx <= 156.2\ngini = 0.366\nclass = 2297\nvalue = [2720, 5, 857]\nclass = a'),  
    Text(3855.272727272727, 543.5999999999999, 'BEN <= 1.905\ngini = 0.489\nclasses = 338\nvalue = [299, 1, 215]\nclass = a'),  
    Text(3753.8181818182, 181.19999999999982, 'gini = 0.185\nclass = 124\nvalue = [19, 0, 165]\nclass = c'),  
    Text(3956.72727272725, 181.19999999999982, 'gini = 0.262\nclass = 214\nvalue = [280, 1, 50]\nclass = a'),  
    Text(4261.090909090909, 543.5999999999999, 'TCH <= 1.485\ngini = 0.333\nclasses = 1959\nvalue = [2421, 4, 642]\nclass = a'),  
    Text(4159.6363636364, 181.19999999999982, 'gini = 0.144\nclass = 740\nvalue = [1077, 0, 91]\nclass = a'),  
    Text(4362.545454545454, 181.19999999999982, 'gini = 0.415\nclass = 1219\nvalue = [1344, 4, 551]\nclass = a')]
```



## Accuracy

# Linear Regression

```
In [63]: lr.score(x_train,y_train)
```

Out[63]: 0.31282390667177273

# Ridge Regression

```
In [64]: rr.score(x_train,y_train)
```

Out[64]: 0.3125741505664318

# Lasso Regression

```
In [65]: la.score(x_test,y_test)
```

Out[65]: 0.06685746388029112

# ElasticNet Regression

```
In [66]: en.score(x_test,y_test)
```

**Out[66]:** 0.17339846972558737

## Logistic Regression

```
In [67]: logr.score(fs,target_vector)
```

```
Out[67]: 0.879023418036871
```

## Random Forest

```
In [68]: grid_search.best_score_
```

```
Out[68]: 0.8633356706847494
```

## Conclusion

**Logistic Regression is suitable for this dataset**