

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [2]: df=pd.read_csv(r"E:\154\C10_air\csvs_per_year\csvs_per_year\madrid_2011.csv")
df
```

Out[2]:

| | date | BEN | CO | EBE | NMHC | NO | NO_2 | O_3 | PM10 | PM25 | SO_2 | TCH | TOL |
|--------|---------------------|-----|-----|-----|------|-------|------|------|------|------|------|------|-------|
| 0 | 2011-11-01 01:00:00 | NaN | 1.0 | NaN | NaN | 154.0 | 84.0 | NaN | NaN | NaN | 6.0 | NaN | NaN 2 |
| 1 | 2011-11-01 01:00:00 | 2.5 | 0.4 | 3.5 | 0.26 | 68.0 | 92.0 | 3.0 | 40.0 | 24.0 | 9.0 | 1.54 | 8.7 2 |
| 2 | 2011-11-01 01:00:00 | 2.9 | NaN | 3.8 | NaN | 96.0 | 99.0 | NaN | NaN | NaN | NaN | NaN | 7.2 2 |
| 3 | 2011-11-01 01:00:00 | NaN | 0.6 | NaN | NaN | 60.0 | 83.0 | 2.0 | NaN | NaN | NaN | NaN | NaN 2 |
| 4 | 2011-11-01 01:00:00 | NaN | NaN | NaN | NaN | 44.0 | 62.0 | 3.0 | NaN | NaN | 3.0 | NaN | NaN 2 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 209923 | 2011-09-01 00:00:00 | NaN | 0.2 | NaN | NaN | 5.0 | 19.0 | 44.0 | NaN | NaN | NaN | NaN | NaN 2 |
| 209924 | 2011-09-01 00:00:00 | NaN | 0.1 | NaN | NaN | 6.0 | 29.0 | NaN | 11.0 | NaN | 7.0 | NaN | NaN 2 |
| 209925 | 2011-09-01 00:00:00 | NaN | NaN | NaN | 0.23 | 1.0 | 21.0 | 28.0 | NaN | NaN | NaN | 1.44 | NaN 2 |
| 209926 | 2011-09-01 00:00:00 | NaN | NaN | NaN | NaN | 3.0 | 15.0 | 48.0 | NaN | NaN | NaN | NaN | NaN 2 |
| 209927 | 2011-09-01 00:00:00 | NaN | NaN | NaN | NaN | 4.0 | 33.0 | 38.0 | 13.0 | NaN | NaN | NaN | NaN 2 |

209928 rows × 14 columns

Data Cleaning and Data Preprocessing

```
In [3]: df=df.dropna()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
   'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16460 entries, 1 to 209910
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      16460 non-null   object 
 1   BEN       16460 non-null   float64
 2   CO        16460 non-null   float64
 3   EBE       16460 non-null   float64
 4   NMHC      16460 non-null   float64
 5   NO        16460 non-null   float64
 6   NO_2      16460 non-null   float64
 7   O_3       16460 non-null   float64
 8   PM10      16460 non-null   float64
 9   PM25      16460 non-null   float64
 10  SO_2      16460 non-null   float64
 11  TCH       16460 non-null   float64
 12  TOL       16460 non-null   float64
 13  station   16460 non-null   int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 1.9+ MB
```

```
In [6]: data=df[['CO' , 'station']]  
data
```

Out[6]:

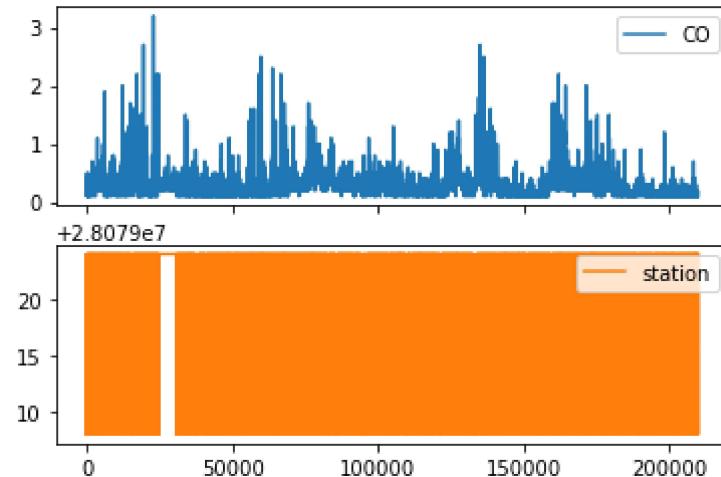
| | CO | station |
|--------|-----|----------|
| 1 | 0.4 | 28079008 |
| 6 | 0.3 | 28079024 |
| 25 | 0.3 | 28079008 |
| 30 | 0.4 | 28079024 |
| 49 | 0.2 | 28079008 |
| ... | ... | ... |
| 209862 | 0.1 | 28079024 |
| 209881 | 0.1 | 28079008 |
| 209886 | 0.1 | 28079024 |
| 209905 | 0.1 | 28079008 |
| 209910 | 0.1 | 28079024 |

16460 rows × 2 columns

Line chart

```
In [7]: data.plot.line(subplots=True)
```

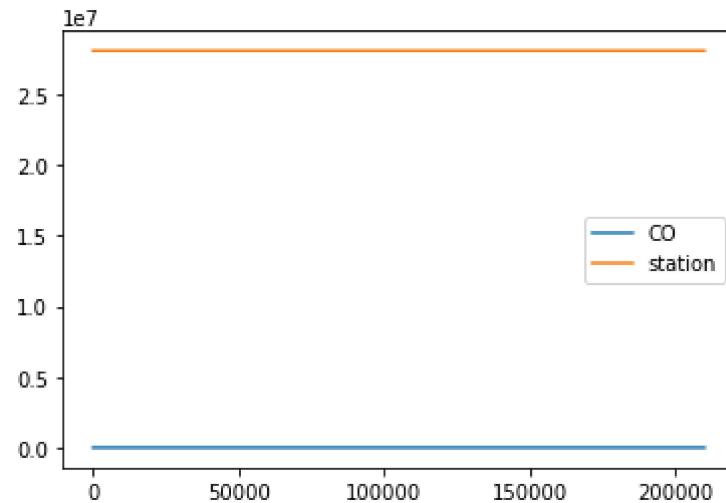
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

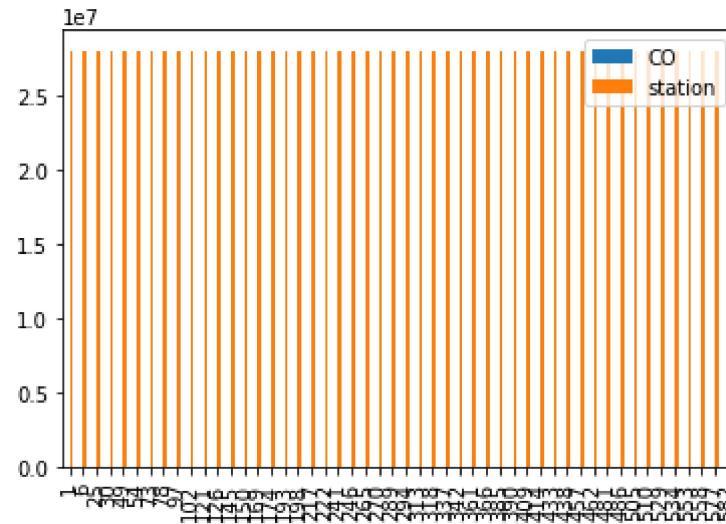


Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

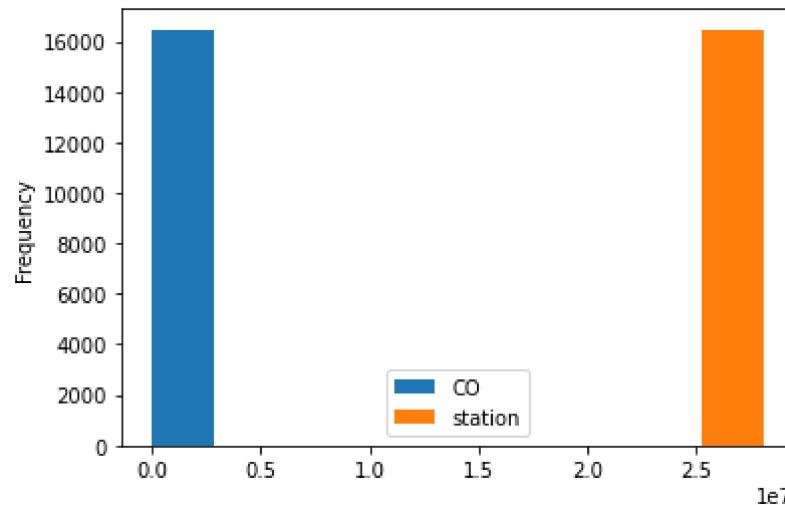
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: data.plot.hist()
```

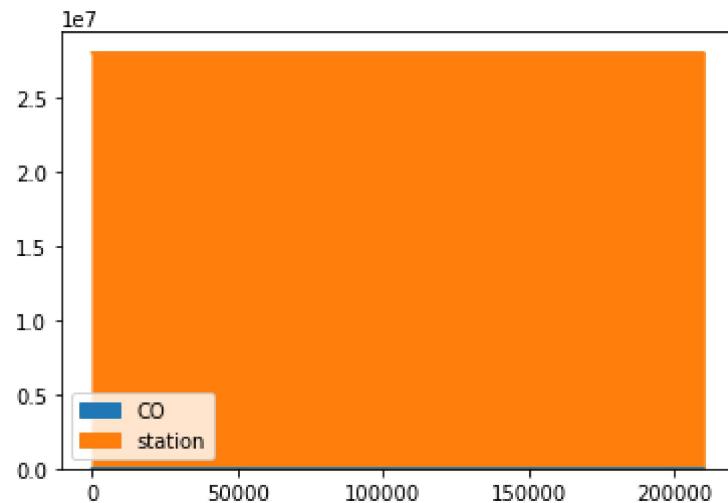
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [12]: data.plot.area()
```

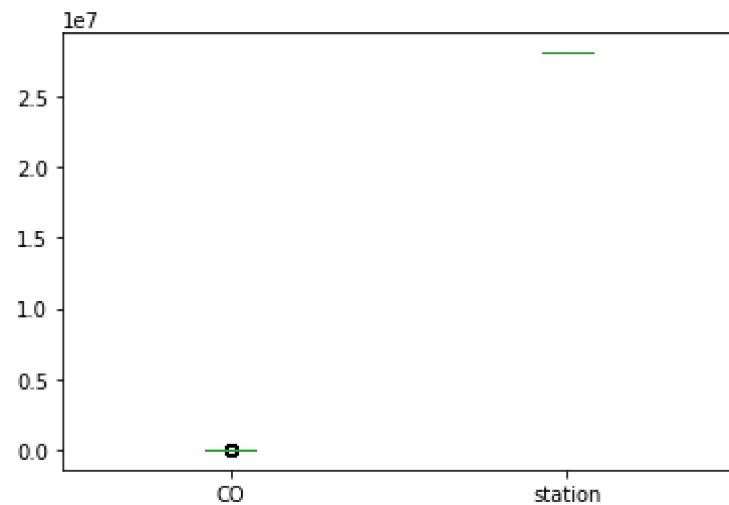
```
Out[12]: <AxesSubplot:>
```



Box chart

In [13]: `data.plot.box()`

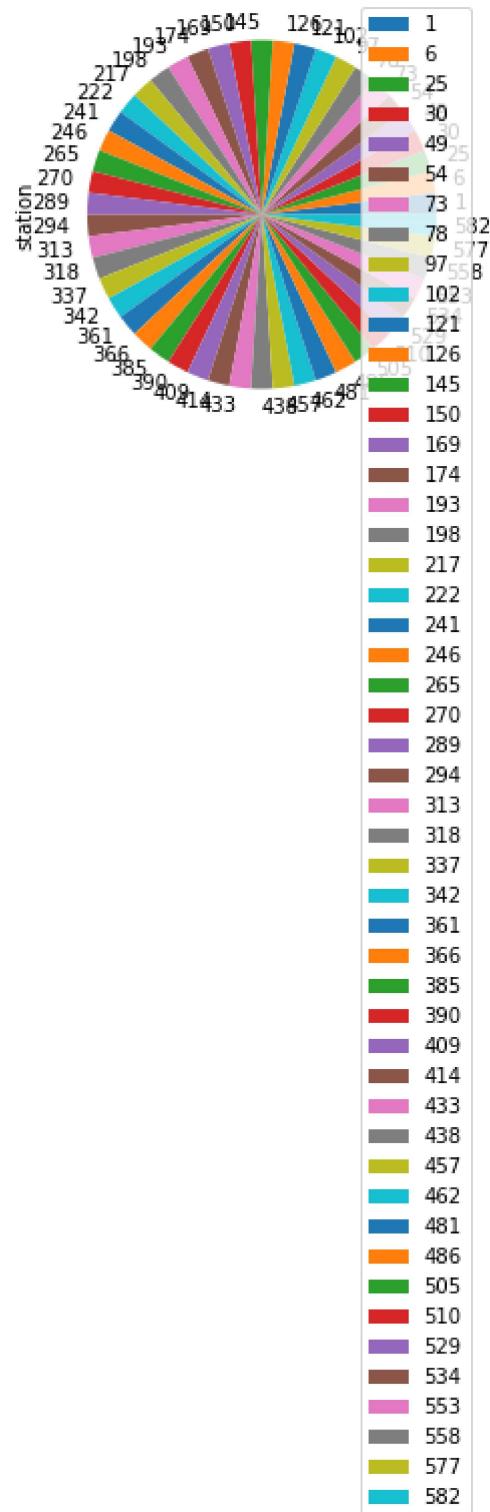
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: b.plot.pie(y='station' )
```

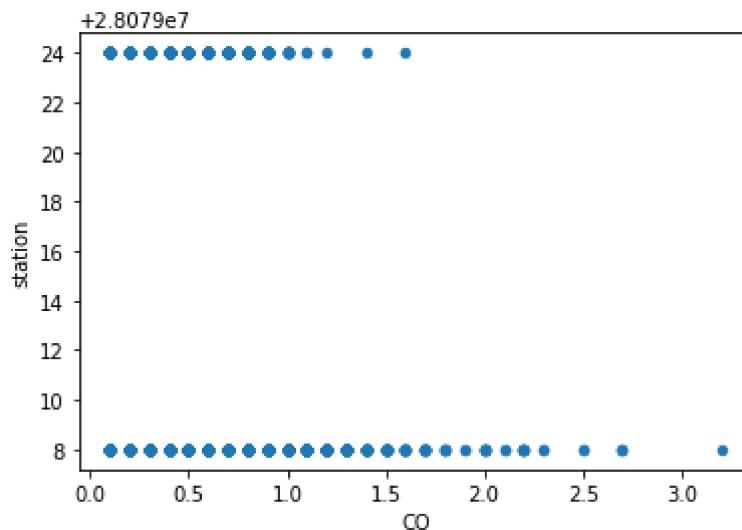
```
Out[14]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16460 entries, 1 to 209910
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      16460 non-null   object 
 1   BEN       16460 non-null   float64
 2   CO        16460 non-null   float64
 3   EBE       16460 non-null   float64
 4   NMHC      16460 non-null   float64
 5   NO        16460 non-null   float64
 6   NO_2      16460 non-null   float64
 7   O_3       16460 non-null   float64
 8   PM10      16460 non-null   float64
 9   PM25      16460 non-null   float64
 10  SO_2      16460 non-null   float64
 11  TCH       16460 non-null   float64
 12  TOL       16460 non-null   float64
 13  station   16460 non-null   int64
```

In [17]: df.describe()

Out[17]:

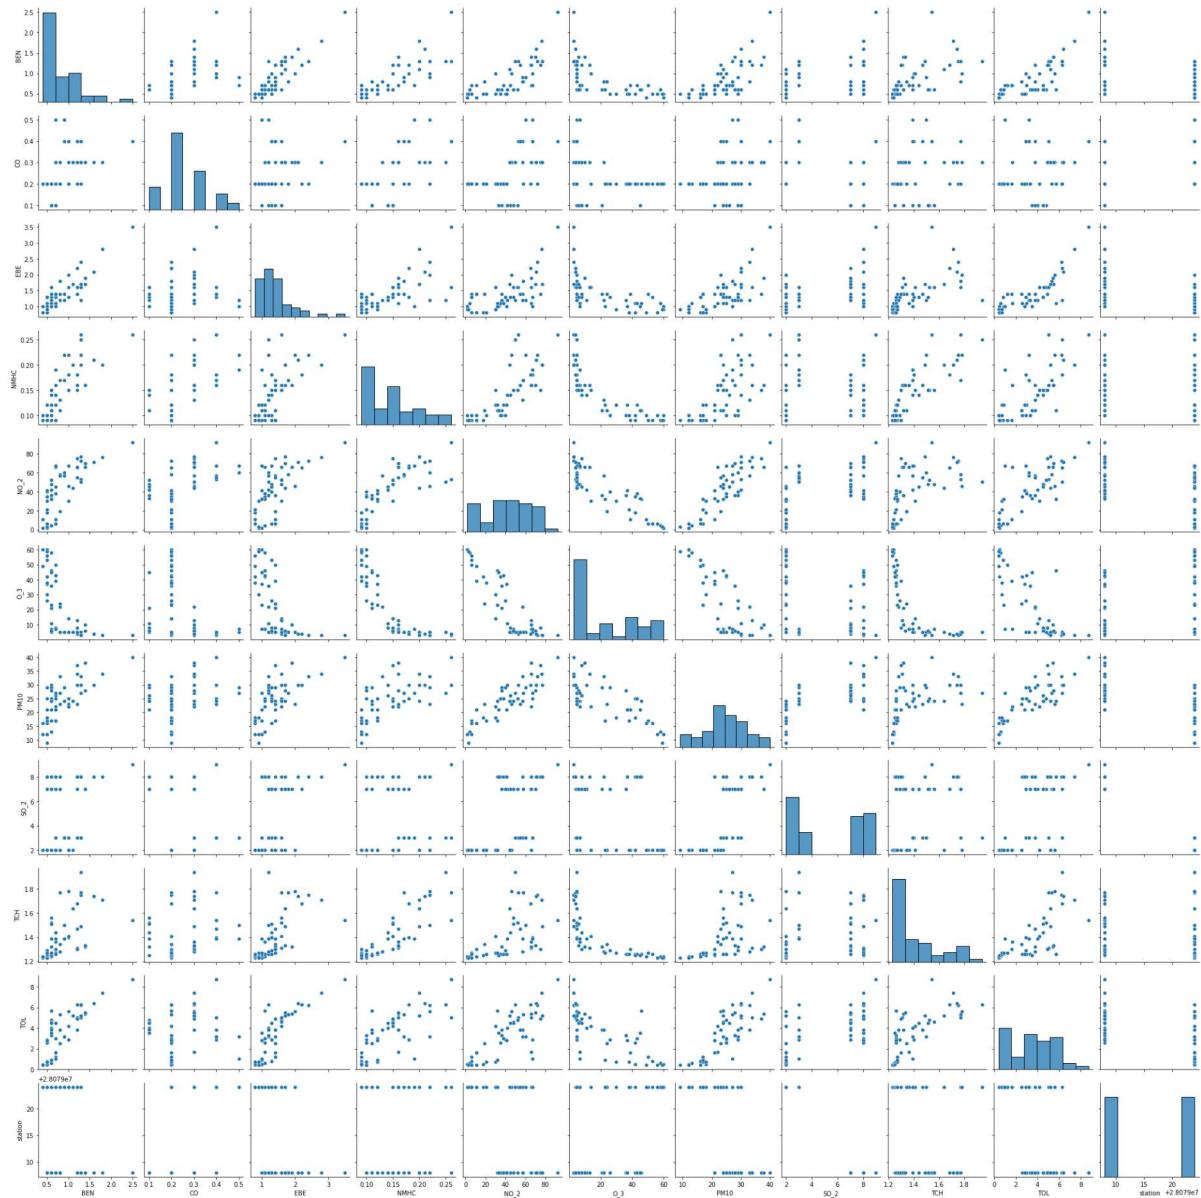
| | BEN | CO | EBE | NMHC | NO | NO_2 | |
|--------------|--------------|--------------|--------------|--------------|--------------|--------------|-----|
| count | 16460.000000 | 16460.000000 | 16460.000000 | 16460.000000 | 16460.000000 | 16460.000000 | 164 |
| mean | 0.900680 | 0.277758 | 1.471871 | 0.167043 | 23.671810 | 44.583961 | |
| std | 0.768892 | 0.206143 | 1.051004 | 0.075068 | 44.362859 | 31.569185 | |
| min | 0.100000 | 0.100000 | 0.200000 | 0.010000 | 1.000000 | 1.000000 | |
| 25% | 0.500000 | 0.200000 | 0.800000 | 0.120000 | 2.000000 | 19.000000 | |
| 50% | 0.700000 | 0.200000 | 1.200000 | 0.160000 | 7.000000 | 40.000000 | |
| 75% | 1.100000 | 0.300000 | 1.700000 | 0.200000 | 25.000000 | 63.000000 | |
| max | 9.500000 | 3.200000 | 12.800000 | 0.840000 | 615.000000 | 289.000000 | 1 |

In [18]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
'PM10', 'SO_2', 'TCH', 'TOL', 'station']]

EDA AND VISUALIZATION

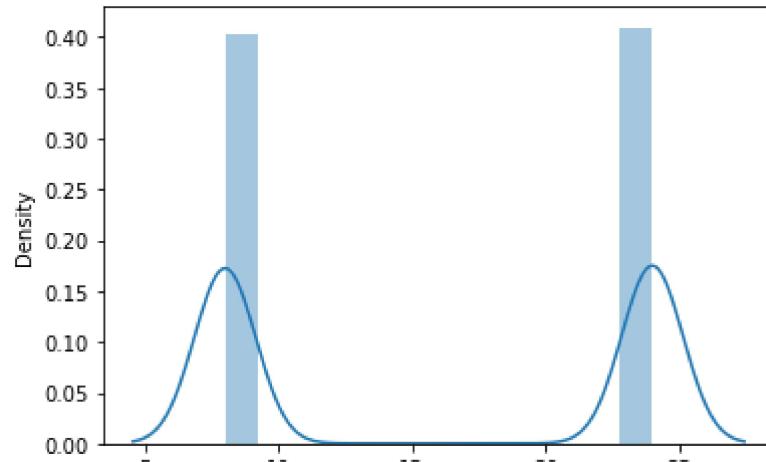
```
In [19]: sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x1b4c1143af0>
```



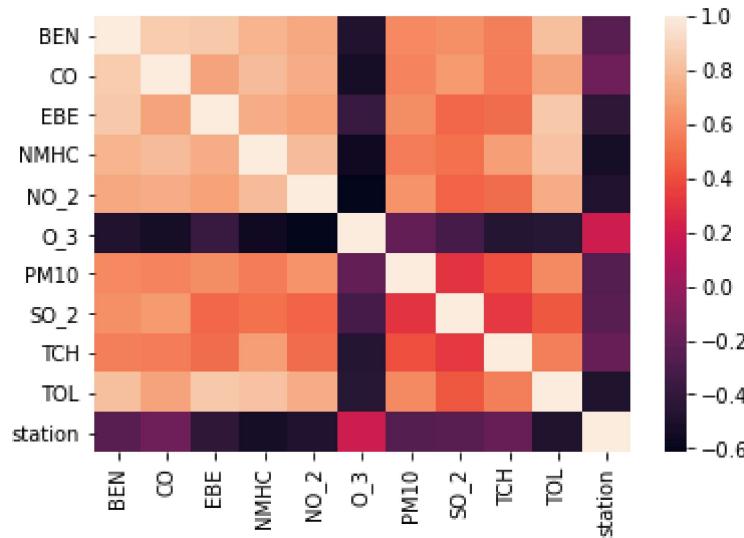
```
In [20]: sns.distplot(df1['station'])
FutureWarning: Please adapt your code to use either `distplot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
           'PM10', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [23]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: 28079018.379894726
```

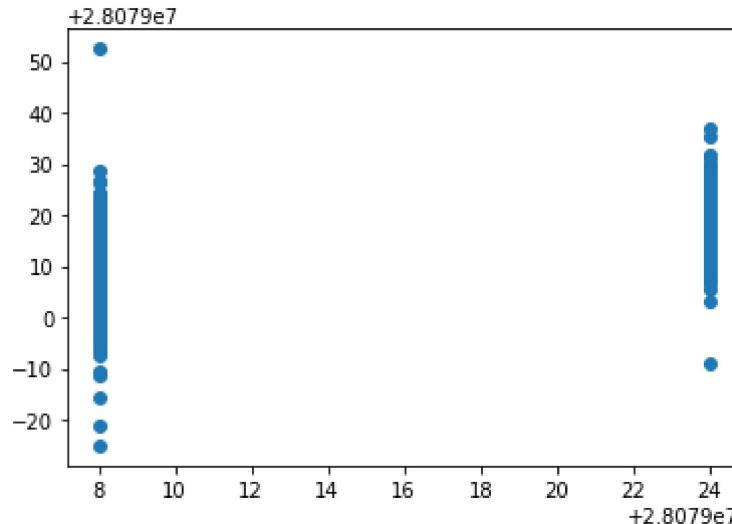
```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[26]:
```

| | Co-efficient |
|------|--------------|
| BEN | 3.957996 |
| CO | 32.795678 |
| EBE | -2.128572 |
| NMHC | -95.653807 |
| NO_2 | -0.081529 |
| O_3 | -0.014847 |
| PM10 | 0.000658 |
| SO_2 | -0.524357 |
| TCH | 9.546366 |
| TOL | -0.447424 |

```
In [27]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1b4c9708ca0>
```



ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.6169329869461591
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.6236818673887985
```

Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.5860635585849482
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.5895649062638599
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

Accuracy(Lasso)

```
In [35]: la.score(x_train,y_train)
```

```
Out[35]: 0.2293062526678089
```

```
In [36]: la.score(x_test,y_test)
```

```
Out[36]: 0.24597881956998158
```

ElasticNet

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: en.coef_
```

```
Out[38]: array([ 6.19490205e-01,  0.00000000e+00, -0.00000000e+00, -0.00000000e+00,
   -1.28854861e-01, -5.36514789e-02,  7.10066717e-02, -4.63639394e-04,
   0.00000000e+00, -7.31046739e-01])
```

```
In [39]: en.intercept_
```

```
Out[39]: 28079024.273564596
```

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

```
Out[41]: 0.3321355459243095
```

Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

5.826072893511698
42.729750322206954
6.536799700327903

Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE','NMHC', 'NO_2','O_3',
                           'PM10', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (16460, 10)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (16460,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10]]
```

```
In [51]: prediction=logr.predict(observation)
print(prediction)
```

[28079008]

```
In [52]: logr.classes_
```

```
Out[52]: array([28079008, 28079024], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.9237545565006076
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.9999999999999966
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[1.0000000e+00, 3.47334507e-15]])
```

Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
}
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                      param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                  'min_samples_leaf': [5, 10, 15, 20, 25],  
                                  'n_estimators': [10, 20, 30, 40, 50]},  
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.9346467627148065
```

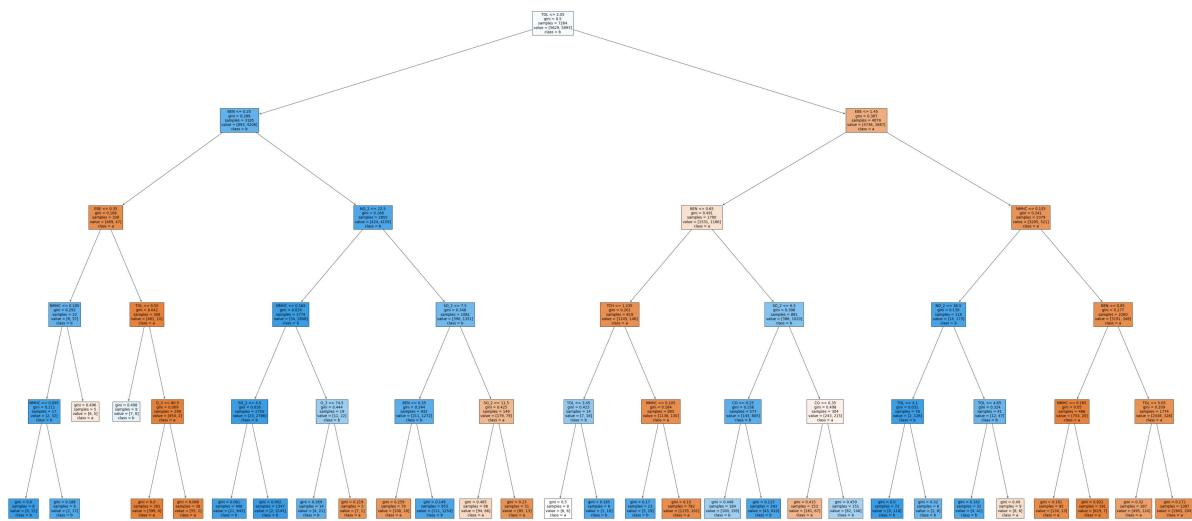
```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b'])
```

```
Out[62]: [Text(2058.8275862068967, 1993.2, 'TOL <= 2.05\nngini = 0.5\nsamples = 7264\nvalue = [5629, 5893]\nnclass = b'),  
Text(885.1034482758621, 1630.8000000000002, 'BEN <= 0.25\nngini = 0.289\nsamples = 3185\nvalue = [893, 4206]\nnclass = b'),  
Text(384.82758620689657, 1268.4, 'EBE <= 0.35\nngini = 0.166\nsamples = 330\nvalue = [469, 47]\nnclass = a'),  
Text(230.89655172413796, 906.0, 'NMHC <= 0.185\nngini = 0.292\nsamples = 22\nvalue = [8, 37]\nnclass = b'),  
Text(153.93103448275863, 543.5999999999999, 'NMHC <= 0.095\nngini = 0.111\nsamples = 17\nvalue = [2, 32]\nnclass = b'),  
Text(76.96551724137932, 181.1999999999982, 'gini = 0.0\nsamples = 8\nvalue = [0, 15]\nnclass = b'),  
Text(230.89655172413796, 181.1999999999982, 'gini = 0.188\nsamples = 9\nvalue = [2, 17]\nnclass = b'),  
Text(307.86206896551727, 543.5999999999999, 'gini = 0.496\nsamples = 5\nvalue = [6, 5]\nnclass = a'),  
Text(538.7586206896552, 906.0, 'TOL <= 0.55\nngini = 0.042\nsamples = 308\nvalue = [461, 10]\nnclass = a'),  
Text(461.79310344827593, 543.5999999999999, 'gini = 0.498\nsamples = 9\nvalue = [7, 8]\nnclass = b'),  
Text(615.7241379310345, 543.5999999999999, 'O_3 <= 80.5\nngini = 0.009\nsamples = 299\nvalue = [454, 2]\nnclass = a'),  
Text(538.7586206896552, 181.1999999999982, 'gini = 0.0\nsamples = 261\nvalue = [399, 0]\nnclass = a'),  
Text(692.6896551724138, 181.1999999999982, 'gini = 0.068\nsamples = 38\nvalue = [55, 2]\nnclass = a'),  
Text(1385.3793103448277, 1268.4, 'NO_2 <= 22.5\nngini = 0.168\nsamples = 2855\nvalue = [424, 4159]\nnclass = b'),  
Text(1077.5172413793105, 906.0, 'NMHC <= 0.165\nngini = 0.024\nsamples = 1774\nvalue = [34, 2808]\nnclass = b'),  
Text(923.5862068965519, 543.5999999999999, 'SO_2 <= 4.5\nngini = 0.016\nsamples = 1755\nvalue = [23, 2786]\nnclass = b'),  
Text(846.6206896551724, 181.1999999999982, 'gini = 0.061\nsamples = 408\nvalue = [21, 643]\nnclass = b'),  
Text(1000.5517241379312, 181.1999999999982, 'gini = 0.002\nsamples = 1347\nvalue = [2, 2143]\nnclass = b'),  
Text(1231.448275862069, 543.5999999999999, 'O_3 <= 74.5\nngini = 0.444\nsamples = 19\nvalue = [11, 22]\nnclass = b'),  
Text(1154.4827586206898, 181.1999999999982, 'gini = 0.269\nsamples = 14\nvalue = [4, 21]\nnclass = b'),  
Text(1308.4137931034484, 181.1999999999982, 'gini = 0.219\nsamples = 5\nvalue = [7, 1]\nnclass = a'),  
Text(1693.2413793103449, 906.0, 'SO_2 <= 7.5\nngini = 0.348\nsamples = 1081\nvalue = [390, 1351]\nnclass = b'),  
Text(1539.3103448275863, 543.5999999999999, 'BEN <= 0.35\nngini = 0.244\nsamples = 932\nvalue = [211, 1272]\nnclass = b'),  
Text(1462.344827586207, 181.1999999999982, 'gini = 0.259\nsamples = 79\nvalue = [100, 18]\nnclass = a'),  
Text(1616.2758620689656, 181.1999999999982, 'gini = 0.149\nsamples = 853\nvalue = [111, 1254]\nnclass = b'),  
Text(1847.1724137931037, 543.5999999999999, 'SO_2 <= 11.5\nngini = 0.425\nsamples = 149\nvalue = [179, 79]\nnclass = a'),  
Text(1770.2068965517242, 181.1999999999982, 'gini = 0.485\nsamples = 98\nvalue = [94, 66]\nnclass = a'),  
Text(1924.137931034483, 181.1999999999982, 'gini = 0.23\nsamples = 51\nvalue = [85, 13]\nnclass = a'),  
Text(3232.551724137931, 1630.8000000000002, 'EBE <= 1.45\nngini = 0.387\nsamples = 17\nvalue = [16, 15]\nnclass = a')]
```

```
les = 4079\nvalue = [4736, 1687]\nclass = a'),  
    Text(2616.8275862068967, 1268.4, 'BEN <= 0.65\nngini = 0.491\nsamples = 1700  
\nvalue = [1531, 1166]\nclass = a'),  
    Text(2308.9655172413795, 906.0, 'TCH <= 1.235\nngini = 0.201\nsamples = 819\n  
value = [1145, 146]\nclass = a'),  
    Text(2155.034482758621, 543.5999999999999, 'TOL <= 3.45\nngini = 0.423\nsampl  
es = 14\nvalue = [7, 16]\nclass = b'),  
    Text(2078.0689655172414, 181.1999999999982, 'gini = 0.5\nsamples = 8\nvalue  
= [6, 6]\nclass = a'),  
    Text(2232.0, 181.1999999999982, 'gini = 0.165\nsamples = 6\nvalue = [1, 10]  
\nclass = b'),  
    Text(2462.896551724138, 543.5999999999999, 'NMHC <= 0.105\nngini = 0.184\nsam  
ples = 805\nvalue = [1138, 130]\nclass = a'),  
    Text(2385.9310344827586, 181.1999999999982, 'gini = 0.17\nsamples = 23\nval  
ue = [3, 29]\nclass = b'),  
    Text(2539.8620689655177, 181.1999999999982, 'gini = 0.15\nsamples = 782\nva  
lue = [1135, 101]\nclass = a'),  
    Text(2924.689655172414, 906.0, 'SO_2 <= 6.5\nngini = 0.398\nsamples = 881\nva  
lue = [386, 1020]\nclass = b'),  
    Text(2770.7586206896553, 543.5999999999999, 'CO <= 0.25\nngini = 0.256\nsampl  
es = 577\nvalue = [143, 805]\nclass = b'),  
    Text(2693.7931034482763, 181.1999999999982, 'gini = 0.448\nsamples = 184\nnv  
alue = [100, 195]\nclass = b'),  
    Text(2847.724137931035, 181.1999999999982, 'gini = 0.123\nsamples = 393\nva  
lue = [43, 610]\nclass = b'),  
    Text(3078.6206896551726, 543.5999999999999, 'CO <= 0.35\nngini = 0.498\nsampl  
es = 304\nvalue = [243, 215]\nclass = a'),  
    Text(3001.6551724137935, 181.1999999999982, 'gini = 0.415\nsamples = 153\nnv  
alue = [161, 67]\nclass = a'),  
    Text(3155.586206896552, 181.1999999999982, 'gini = 0.459\nsamples = 151\nva  
lue = [82, 148]\nclass = b'),  
    Text(3848.275862068966, 1268.4, 'NMHC <= 0.135\nngini = 0.241\nsamples = 2379  
\nvalue = [3205, 521]\nclass = a'),  
    Text(3540.4137931034484, 906.0, 'NO_2 <= 36.5\nngini = 0.139\nsamples = 119\nn  
value = [14, 173]\nclass = b'),  
    Text(3386.4827586206898, 543.5999999999999, 'TOL <= 4.1\nngini = 0.031\nsampl  
es = 78\nvalue = [2, 126]\nclass = b'),  
    Text(3309.5172413793107, 181.1999999999982, 'gini = 0.0\nsamples = 72\nvalu  
e = [0, 118]\nclass = b'),  
    Text(3463.4482758620693, 181.1999999999982, 'gini = 0.32\nsamples = 6\nvalu  
e = [2, 8]\nclass = b'),  
    Text(3694.3448275862074, 543.5999999999999, 'TOL <= 4.65\nngini = 0.324\nsam  
ples = 41\nvalue = [12, 47]\nclass = b'),  
    Text(3617.379310344828, 181.1999999999982, 'gini = 0.162\nsamples = 32\nval  
ue = [4, 41]\nclass = b'),  
    Text(3771.3103448275865, 181.1999999999982, 'gini = 0.49\nsamples = 9\nval  
ue = [8, 6]\nclass = a'),  
    Text(4156.137931034483, 906.0, 'BEN <= 0.85\nngini = 0.177\nsamples = 2260\nnv  
alue = [3191, 348]\nclass = a'),  
    Text(4002.2068965517246, 543.5999999999999, 'NMHC <= 0.165\nngini = 0.05\nsam  
ples = 486\nvalue = [753, 20]\nclass = a'),  
    Text(3925.241379310345, 181.1999999999982, 'gini = 0.161\nsamples = 95\nval  
ue = [134, 13]\nclass = a'),  
    Text(4079.1724137931037, 181.1999999999982, 'gini = 0.022\nsamples = 391\nnv  
alue = [619, 7]\nclass = a'),  
    Text(4310.068965517242, 543.5999999999999, 'TOL <= 5.05\nngini = 0.209\nsampl  
es = 1774\nvalue = [2438, 328]\nclass = a'),
```

```
Text(4233.103448275862, 181.19999999999982, 'gini = 0.32\nsamples = 387\nvalue = [495, 124]\nclass = a'),  
Text(4387.034482758621, 181.19999999999982, 'gini = 0.172\nsamples = 1387\nvalue = [1943, 204]\nclass = a')]
```



Accuracy

Linear Regression

```
In [63]: lr.score(x_train,y_train)
```

```
Out[63]: 0.6236818673887985
```

Ridge Regression

```
In [64]: rr.score(x_train,y_train)
```

```
Out[64]: 0.5895649062638599
```

Lasso Regression

```
In [65]: la.score(x_test,y_test)
```

```
Out[65]: 0.24597881956998158
```

ElasticNet Regression

```
In [66]: en.score(x_test,y_test)
```

```
Out[66]: 0.3321355459243095
```

Logistic Regression

```
In [67]: logr.score(fs,target_vector)
```

```
Out[67]: 0.9237545565006076
```

Random Forest

```
In [68]: grid_search.best_score_
```

```
Out[68]: 0.9346467627148065
```

Conclusion

RandomForest for this dataset