

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [2]: df=pd.read_csv(r"E:\154\C10_air\csvs_per_year\csvs_per_year\madrid_2012.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
0	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	7.0	18.0	NaN	NaN	NaN	2.0	NaN	NaN
1	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4
2	2012-09-01 01:00:00	0.4	NaN	0.7	NaN	2.0	10.0	NaN	NaN	NaN	NaN	NaN	1.5
3	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	50.0	NaN	NaN	NaN	NaN	2.8
4	2012-09-01 01:00:00	NaN	NaN	NaN	NaN	1.0	13.0	54.0	NaN	NaN	3.0	NaN	NaN
...
210715	2012-03-01 00:00:00	NaN	0.6	NaN	NaN	37.0	84.0	14.0	NaN	NaN	NaN	NaN	2.8
210716	2012-03-01 00:00:00	NaN	0.4	NaN	NaN	5.0	76.0	NaN	17.0	NaN	7.0	NaN	NaN
210717	2012-03-01 00:00:00	NaN	NaN	NaN	0.34	3.0	41.0	24.0	NaN	NaN	NaN	1.34	NaN
210718	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	44.0	36.0	NaN	NaN	NaN	NaN	2.8
210719	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	56.0	40.0	18.0	NaN	NaN	NaN	2.8

210720 rows × 14 columns

Data Cleaning and Data Preprocessing

```
In [3]: df=df.dropna()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25',
       'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10916 entries, 6 to 210702
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      10916 non-null   object 
 1   BEN        10916 non-null   float64
 2   CO         10916 non-null   float64
 3   EBE        10916 non-null   float64
 4   NMHC       10916 non-null   float64
 5   NO         10916 non-null   float64
 6   NO_2       10916 non-null   float64
 7   O_3        10916 non-null   float64
 8   PM10       10916 non-null   float64
 9   PM25       10916 non-null   float64
 10  SO_2       10916 non-null   float64
 11  TCH        10916 non-null   float64
 12  TOL        10916 non-null   float64
 13  station    10916 non-null   int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 1.2+ MB
```

In [6]: `data=df[['CO' , 'station']]
data`

Out[6]:

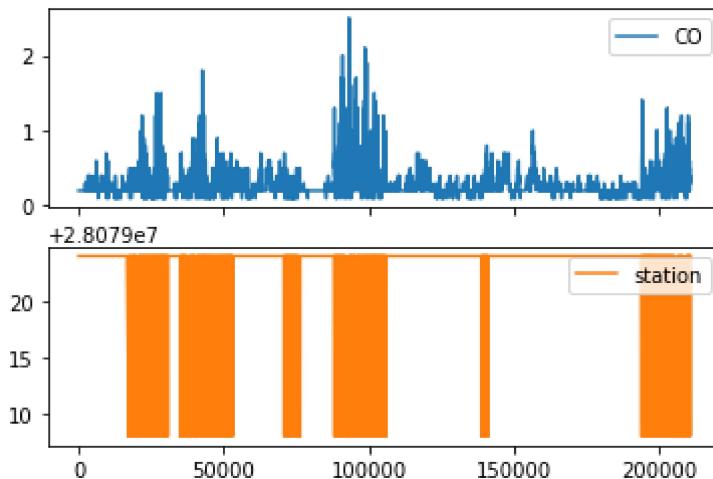
	CO	station
6	0.2	28079024
30	0.2	28079024
54	0.2	28079024
78	0.2	28079024
102	0.2	28079024
...
210654	0.3	28079024
210673	0.4	28079008
210678	0.3	28079024
210697	0.4	28079008
210702	0.3	28079024

10916 rows × 2 columns

Line chart

In [7]: `data.plot.line(subplots=True)`

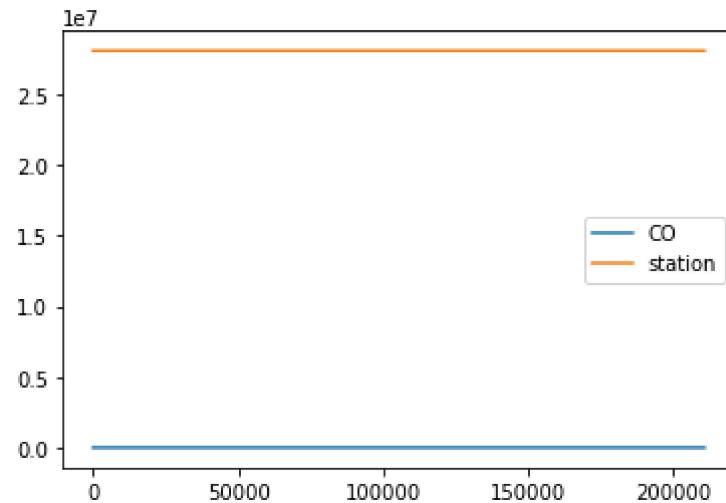
Out[7]: `array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)`



Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

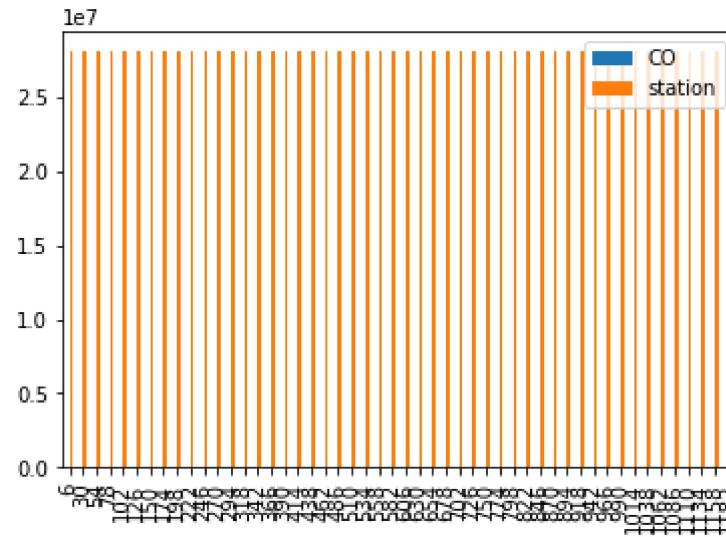


Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

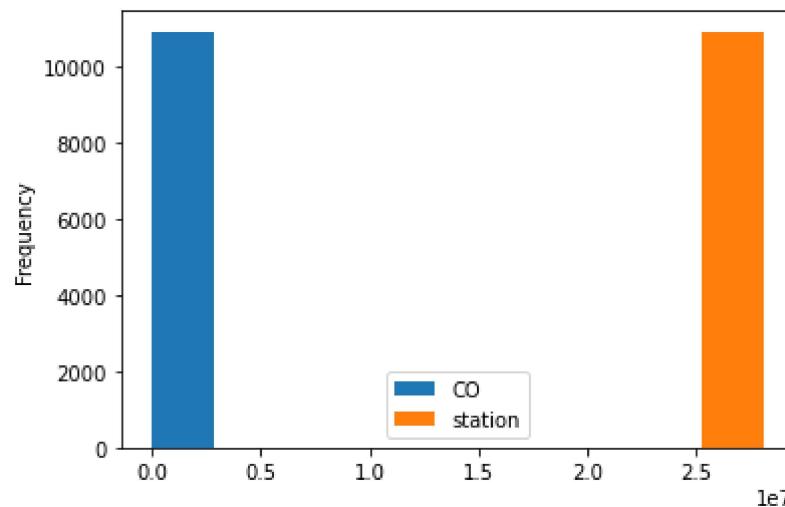
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: data.plot.hist()
```

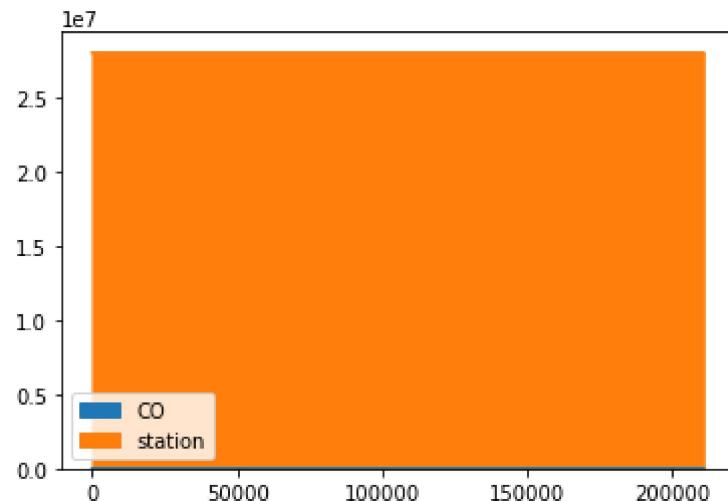
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [12]: data.plot.area()
```

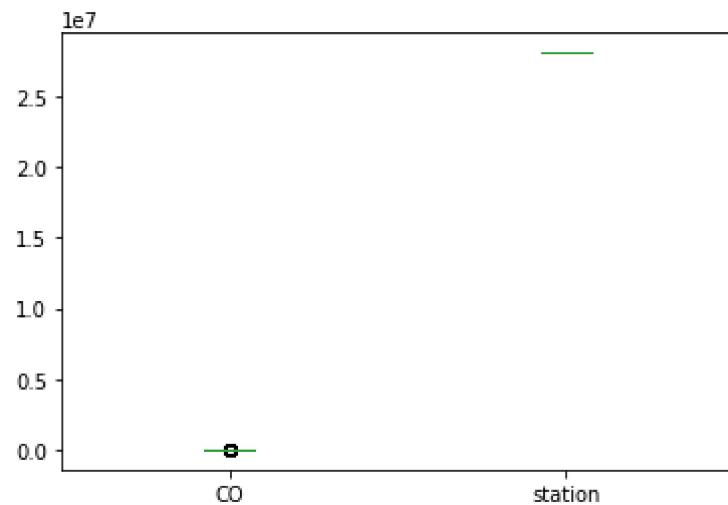
```
Out[12]: <AxesSubplot:>
```



Box chart

In [13]: `data.plot.box()`

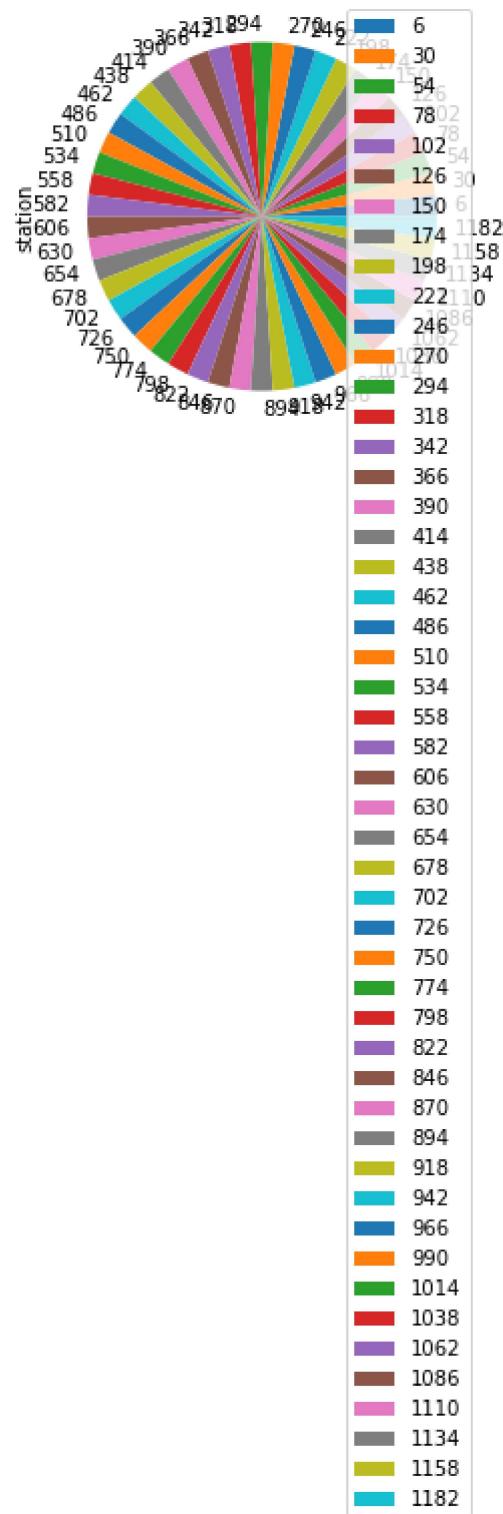
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: b.plot.pie(y='station' )
```

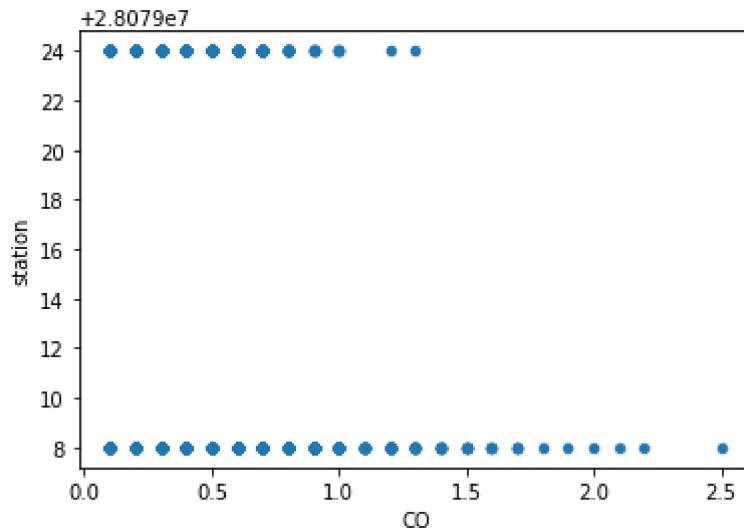
```
Out[14]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: df.info()
```

```
Introducing: 10916 entries, 0 to 210702
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype  
 --- 
 0   date      10916 non-null   object 
 1   BEN       10916 non-null   float64
 2   CO        10916 non-null   float64
 3   EBE       10916 non-null   float64
 4   NMHC      10916 non-null   float64
 5   NO        10916 non-null   float64
 6   NO_2      10916 non-null   float64
 7   O_3       10916 non-null   float64
 8   PM10      10916 non-null   float64
 9   PM25      10916 non-null   float64
 10  SO_2      10916 non-null   float64
 11  TCH       10916 non-null   float64
 12  TOL       10916 non-null   float64
 13  station   10916 non-null   int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 1.2+ MB
```

In [17]: df.describe()

Out[17]:

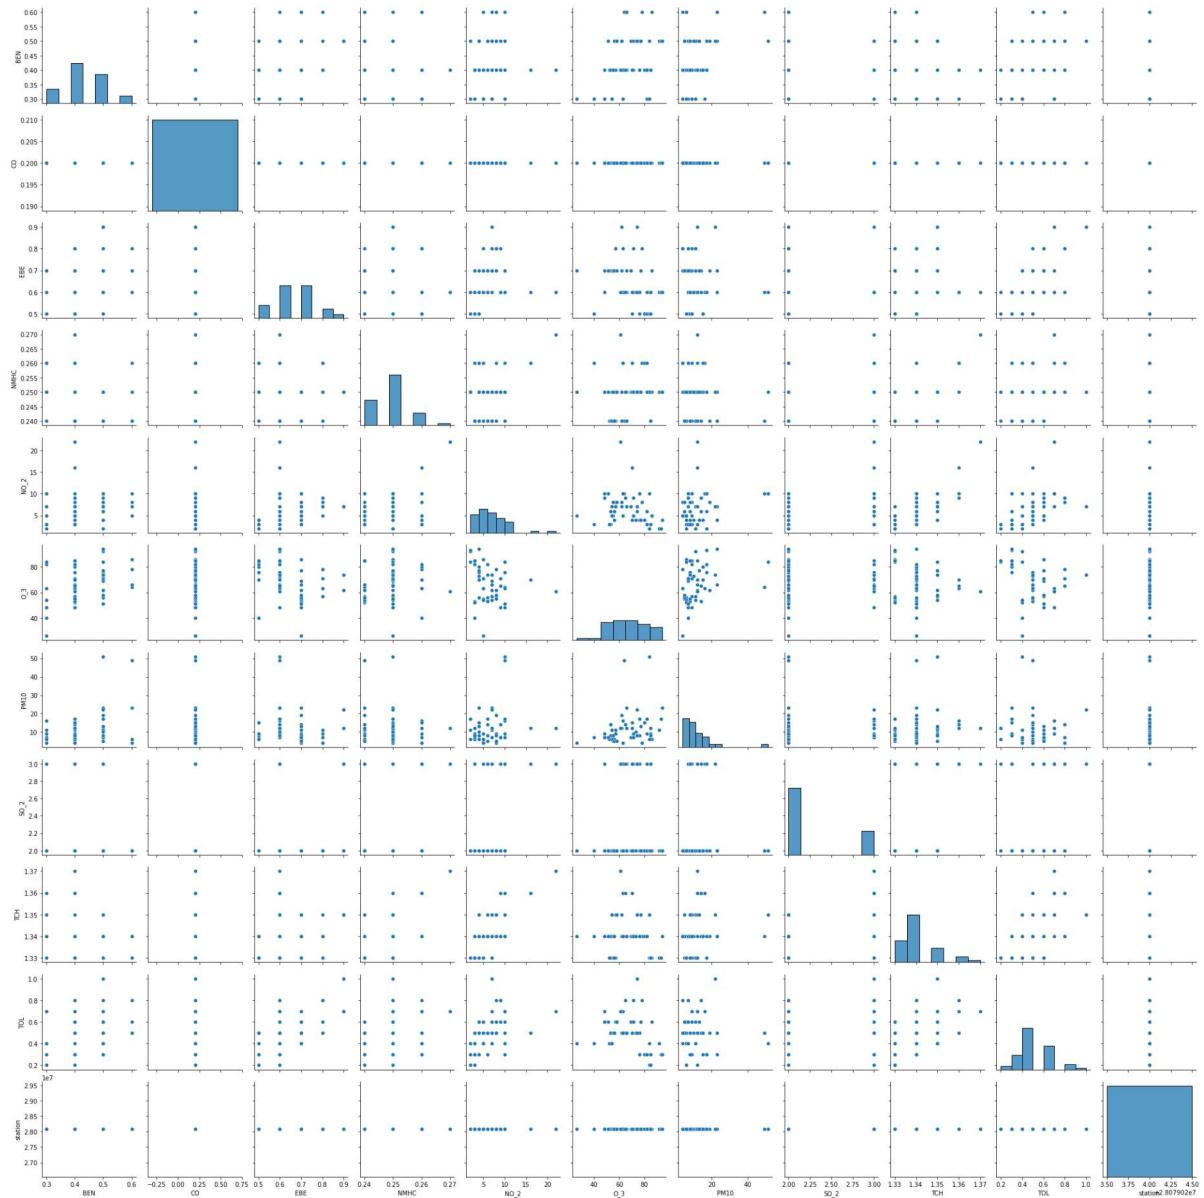
	BEN	CO	EBE	NMHC	NO	NO_2	
count	10916.000000	10916.000000	10916.000000	10916.000000	10916.000000	10916.000000	109
mean	0.784014	0.279333	0.992213	0.215755	18.795529	31.262642	
std	0.632755	0.167922	0.804554	0.075169	40.038872	27.234732	
min	0.100000	0.100000	0.100000	0.050000	0.000000	1.000000	
25%	0.400000	0.200000	0.500000	0.160000	1.000000	9.000000	
50%	0.600000	0.200000	0.800000	0.220000	3.000000	24.000000	
75%	0.900000	0.300000	1.200000	0.250000	18.000000	47.000000	
max	7.000000	2.500000	9.700000	0.670000	525.000000	225.000000	1

In [18]: df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
'PM10', 'SO_2', 'TCH', 'TOL', 'station']]

EDA AND VISUALIZATION

```
In [19]: sns.pairplot(df1[0:50])
```

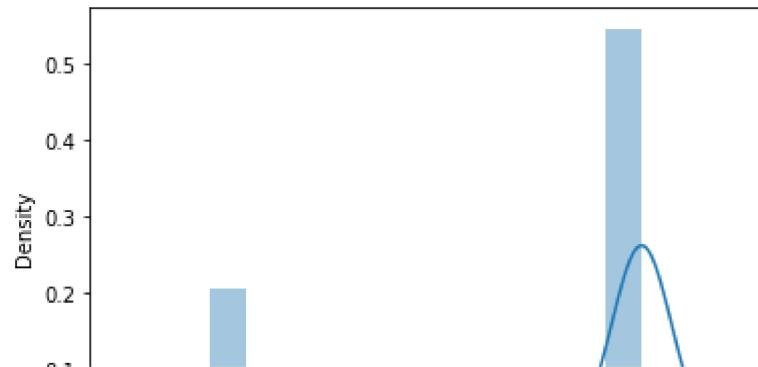
```
Out[19]: <seaborn.axisgrid.PairGrid at 0x29fb5d7d430>
```



In [20]: `sns.distplot(df1['station'])`

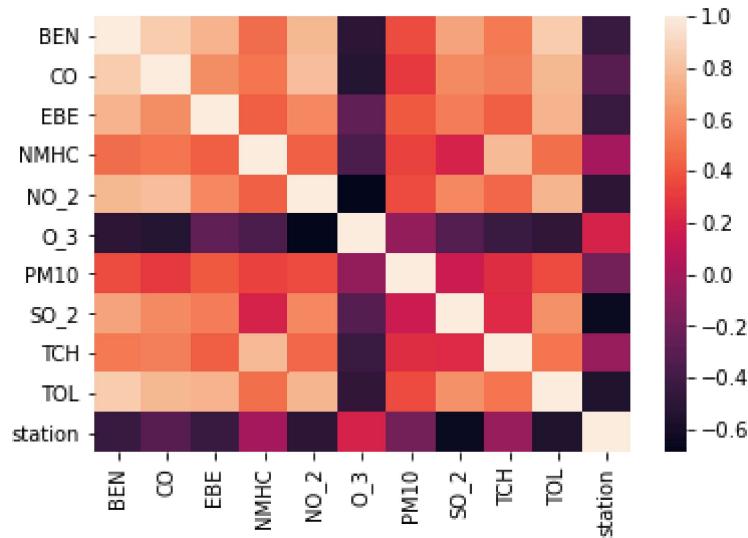
```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

In [22]: `x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3', 'PM10', 'SO_2', 'TCH', 'TOL']]
y=df['station']`

```
In [23]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: 28079019.80932486
```

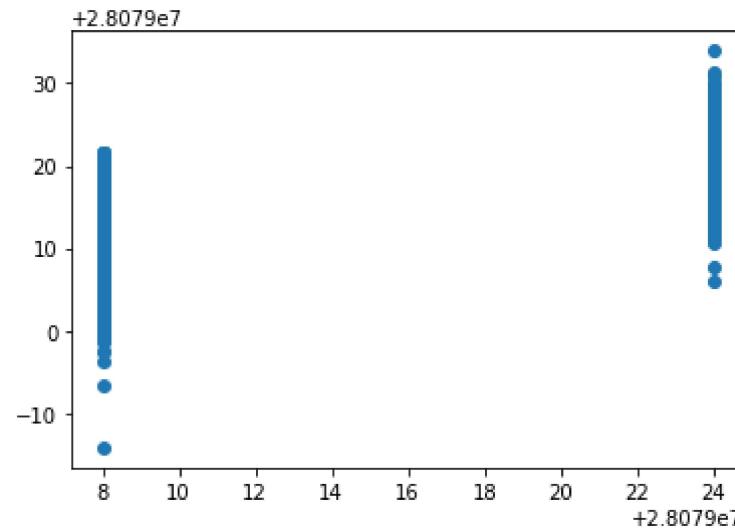
```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[26]:
```

	Co-efficient
BEN	3.236885
CO	17.411361
EBE	-0.502580
NMHC	17.216789
NO_2	-0.116675
O_3	-0.028662
PM10	-0.011613
SO_2	-0.710634
TCH	0.651984
TOL	-1.304644

```
In [27]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x29fbf0ff520>
```



ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.6394723610807198
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.6130715671849114
```

Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.6365049510579274
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.6105151639688398
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

Accuracy(Lasso)

```
In [35]: la.score(x_train,y_train)
```

```
Out[35]: 0.36826562637142735
```

```
In [36]: la.score(x_test,y_test)
```

```
Out[36]: 0.3687663704003382
```

ElasticNet

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: en.coef_
```

```
Out[38]: array([ 0.          ,  0.          , -0.          ,  0.          ,
 -0.03794886,  0.          , -0.60477444,  0.          ,
 -0.28845714])
```

```
In [39]: en.intercept_
```

```
Out[39]: 28079026.719892986
```

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

```
Out[41]: 0.46918672554636864
```

Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

3.7760161348619987
27.266642819000104
5.221747104083087

Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE','NMHC', 'NO_2','O_3',
                           'PM10', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (10916, 10)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (10916,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10]]
```

```
In [51]: prediction=logr.predict(observation)
print(prediction)
```

[28079008]

```
In [52]: logr.classes_
```

```
Out[52]: array([28079008, 28079024], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.9293697325027482
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 1.0
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[1.0, 3.50349553e-26]])
```

Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
}
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                      param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                  'min_samples_leaf': [5, 10, 15, 20, 25],  
                                  'n_estimators': [10, 20, 30, 40, 50]},  
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

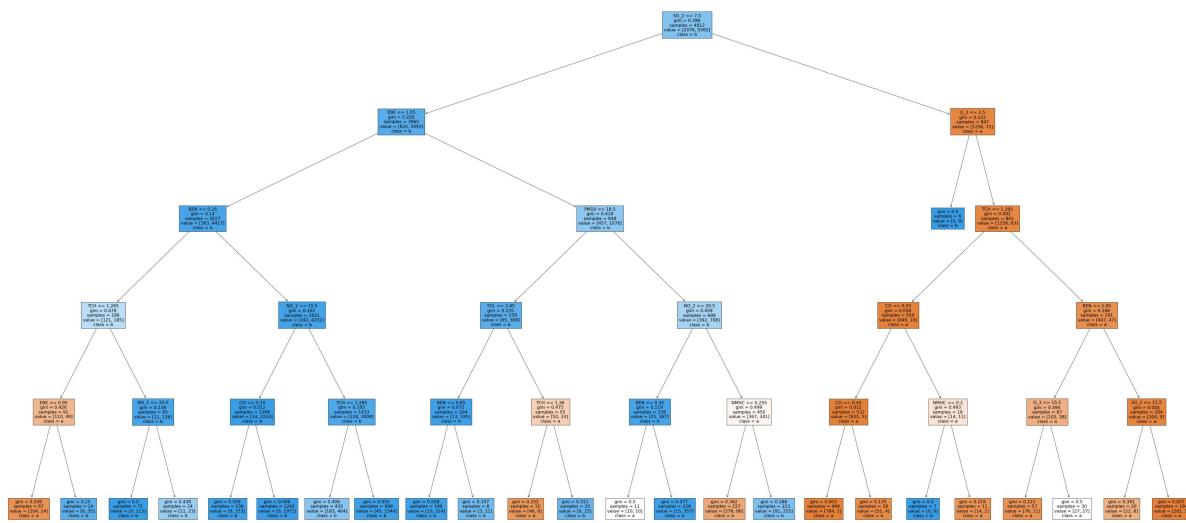
```
Out[60]: 0.9592984348002428
```

```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b'])
```

```
Out[62]: [Text(2557.5, 1993.2, 'SO_2 <= 7.5\ngini = 0.396\nsamples = 4812\nvalue = [20  
76, 5565]\nclass = b'),  
Text(1488.0, 1630.800000000002, 'EBE <= 1.05\ngini = 0.226\nsamples = 3965  
\nvalue = [820, 5493]\nclass = b'),  
Text(744.0, 1268.4, 'BEN <= 0.25\ngini = 0.14\nsamples = 3017\nvalue = [363,  
4417]\nclass = b'),  
Text(372.0, 906.0, 'TCH <= 1.265\ngini = 0.478\nsamples = 186\nvalue = [121,  
185]\nclass = b'),  
Text(186.0, 543.599999999999, 'EBE <= 0.95\ngini = 0.426\nsamples = 91\nval  
ue = [110, 49]\nclass = a'),  
Text(93.0, 181.1999999999982, 'gini = 0.209\nsamples = 67\nvalue = [104, 1  
4]\nclass = a'),  
Text(279.0, 181.1999999999982, 'gini = 0.25\nsamples = 24\nvalue = [6, 35]  
\nclass = b'),  
Text(558.0, 543.599999999999, 'NO_2 <= 20.0\ngini = 0.138\nsamples = 95\nva  
lue = [11, 136]\nclass = b'),  
Text(465.0, 181.1999999999982, 'gini = 0.0\ngsamples = 71\nvalue = [0, 113]  
\nclass = b'),  
Text(651.0, 181.1999999999982, 'gini = 0.438\nsamples = 24\nvalue = [11, 2  
3]\nclass = b'),  
Text(1116.0, 906.0, 'NO_2 <= 15.5\ngini = 0.102\nsamples = 2831\nvalue = [24  
2, 4232]\nclass = b'),  
Text(930.0, 543.599999999999, 'CO <= 0.15\ngini = 0.012\nsamples = 1398\nva  
lue = [14, 2224]\nclass = b'),  
Text(837.0, 181.1999999999982, 'gini = 0.059\nsamples = 156\nvalue = [8, 25  
3]\nclass = b'),  
Text(1023.0, 181.1999999999982, 'gini = 0.006\nsamples = 1242\nvalue = [6,  
1971]\nclass = b'),  
Text(1302.0, 543.599999999999, 'TCH <= 1.285\ngini = 0.183\nsamples = 1433  
\nvalue = [228, 2008]\nclass = b'),  
Text(1209.0, 181.1999999999982, 'gini = 0.406\nsamples = 435\nvalue = [183,  
464]\nclass = b'),  
Text(1395.0, 181.1999999999982, 'gini = 0.055\nsamples = 998\nvalue = [45,  
1544]\nclass = b'),  
Text(2232.0, 1268.4, 'PM10 <= 18.5\ngini = 0.418\nsamples = 948\nvalue = [45  
7, 1076]\nclass = b'),  
Text(1860.0, 906.0, 'TOL <= 2.85\ngini = 0.255\nsamples = 259\nvalue = [65,  
368]\nclass = b'),  
Text(1674.0, 543.599999999999, 'BEN <= 0.85\ngini = 0.072\nsamples = 204\nva  
lue = [13, 335]\nclass = b'),  
Text(1581.0, 181.1999999999982, 'gini = 0.058\nsamples = 196\nvalue = [10,  
324]\nclass = b'),  
Text(1767.0, 181.1999999999982, 'gini = 0.337\nsamples = 8\nvalue = [3, 11]  
\nclass = b'),  
Text(2046.0, 543.599999999999, 'TCH <= 1.36\ngini = 0.475\nsamples = 55\nva  
lue = [52, 33]\nclass = a'),  
Text(1953.0, 181.1999999999982, 'gini = 0.252\nsamples = 35\nvalue = [46,  
8]\nclass = a'),  
Text(2139.0, 181.1999999999982, 'gini = 0.312\nsamples = 20\nvalue = [6, 2  
5]\nclass = b'),  
Text(2604.0, 906.0, 'NO_2 <= 29.5\ngini = 0.459\nsamples = 689\nvalue = [39  
2, 708]\nclass = b'),  
Text(2418.0, 543.599999999999, 'BEN <= 0.35\ngini = 0.119\nsamples = 239\nva  
lue = [25, 367]\nclass = b'),  
Text(2325.0, 181.1999999999982, 'gini = 0.5\nsamples = 11\nvalue = [10, 10]  
\nclass = a'),  
Text(2511.0, 181.1999999999982, 'gini = 0.077\nsamples = 228\nvalue = [15,
```

```
357]\nclass = b'),
Text(2790.0, 543.599999999999, 'NMHC <= 0.255\ngini = 0.499\nsamples = 450
\nvalue = [367, 341]\nclass = a'),
Text(2697.0, 181.1999999999982, 'gini = 0.362\nsamples = 227\nvalue = [276,
86]\nclass = a'),
Text(2883.0, 181.1999999999982, 'gini = 0.388\nsamples = 223\nvalue = [91,
255]\nclass = b'),
Text(3627.0, 1630.800000000002, 'O_3 <= 2.5\ngini = 0.103\nsamples = 847\nv
alue = [1256, 72]\nclass = a'),
Text(3534.0, 1268.4, 'gini = 0.0\nsamples = 6\nvalue = [0, 9]\nclass = b'),
Text(3720.0, 1268.4, 'TCH <= 1.365\ngini = 0.091\nsamples = 841\nvalue = [12
56, 63]\nclass = a'),
Text(3348.0, 906.0, 'CO <= 0.55\ngini = 0.036\nsamples = 550\nvalue = [849,
16]\nclass = a'),
Text(3162.0, 543.599999999999, 'CO <= 0.45\ngini = 0.012\nsamples = 532\nva
lue = [835, 5]\nclass = a'),
Text(3069.0, 181.1999999999982, 'gini = 0.003\nsamples = 494\nvalue = [784,
1]\nclass = a'),
Text(3255.0, 181.1999999999982, 'gini = 0.135\nsamples = 38\nvalue = [51,
4]\nclass = a'),
Text(3534.0, 543.599999999999, 'NMHC <= 0.2\ngini = 0.493\nsamples = 18\nc
lue = [14, 11]\nclass = a'),
Text(3441.0, 181.1999999999982, 'gini = 0.0\nsamples = 7\nvalue = [0, 9]\nc
lass = b'),
Text(3627.0, 181.1999999999982, 'gini = 0.219\nsamples = 11\nvalue = [14,
2]\nclass = a'),
Text(4092.0, 906.0, 'BEN <= 1.95\ngini = 0.186\nsamples = 291\nvalue = [407,
47]\nclass = a'),
Text(3906.0, 543.599999999999, 'O_3 <= 15.5\ngini = 0.394\nsamples = 87\nc
lue = [103, 38]\nclass = a'),
Text(3813.0, 181.1999999999982, 'gini = 0.221\nsamples = 57\nvalue = [76, 1
1]\nclass = a'),
Text(3999.0, 181.1999999999982, 'gini = 0.5\nsamples = 30\nvalue = [27, 27]
\nclass = a'),
Text(4278.0, 543.599999999999, 'SO_2 <= 11.5\ngini = 0.056\nsamples = 204\nc
lue = [304, 9]\nclass = a'),
Text(4185.0, 181.1999999999982, 'gini = 0.391\nsamples = 20\nvalue = [22,
8]\nclass = a'),
Text(4371.0, 181.1999999999982, 'gini = 0.007\nsamples = 184\nvalue = [282,
1]\nclass = a')]
```



Accuracy

Linear Regression

```
In [63]: lr.score(x_train,y_train)
```

```
Out[63]: 0.6130715671849114
```

Ridge Regression

```
In [64]: rr.score(x_train,y_train)
```

```
Out[64]: 0.6105151639688398
```

Lasso Regression

```
In [65]: la.score(x_test,y_test)
```

```
Out[65]: 0.3687663704003382
```

ElasticNet Regression

```
In [66]: en.score(x_test,y_test)
```

```
Out[66]: 0.46918672554636864
```

Logistic Regression

```
In [67]: logr.score(fs,target_vector)
```

```
Out[67]: 0.9293697325027482
```

Random Forest

```
In [68]: grid_search.best_score_
```

```
Out[68]: 0.9592984348002428
```

Conclusion

RandomForest for this dataset